



# Docker – Day 4

Kubernetes

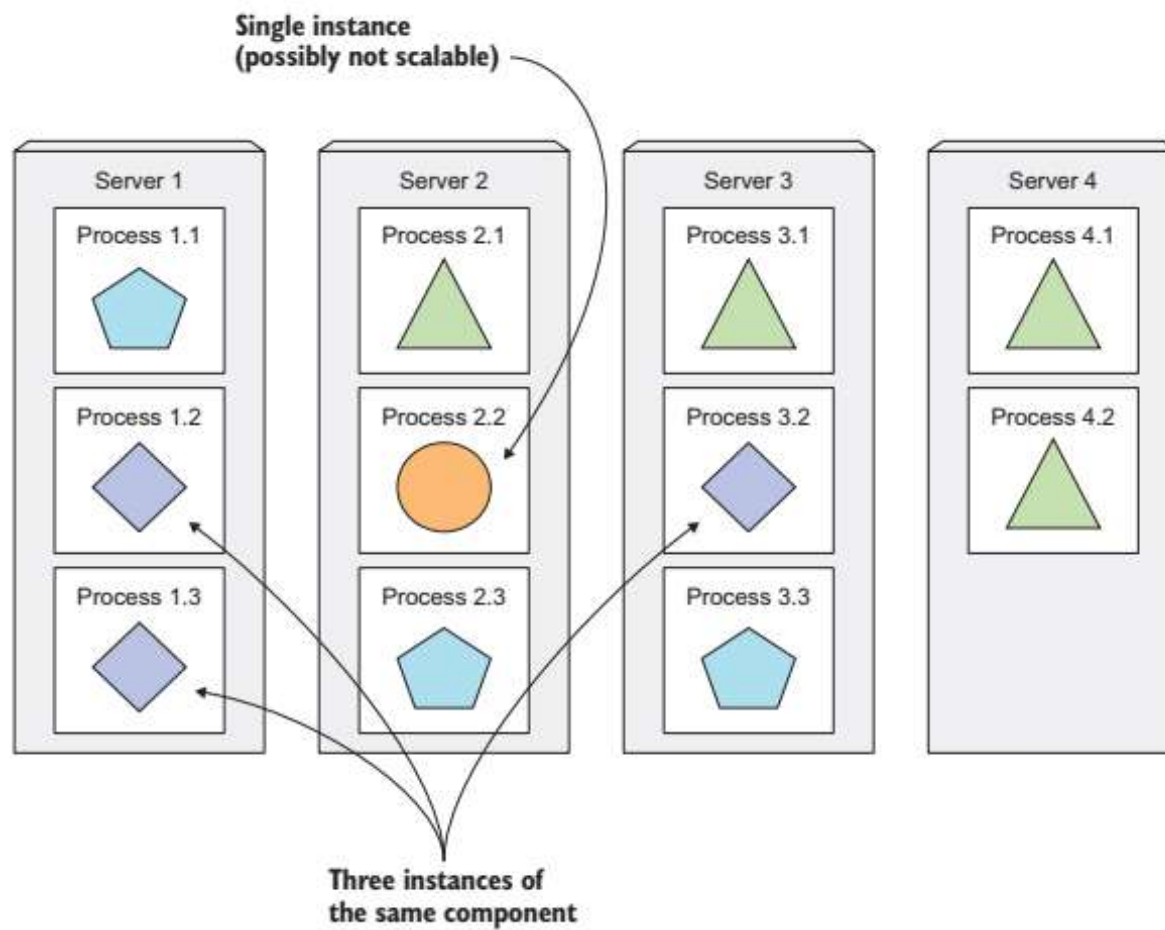
## Monolithic application

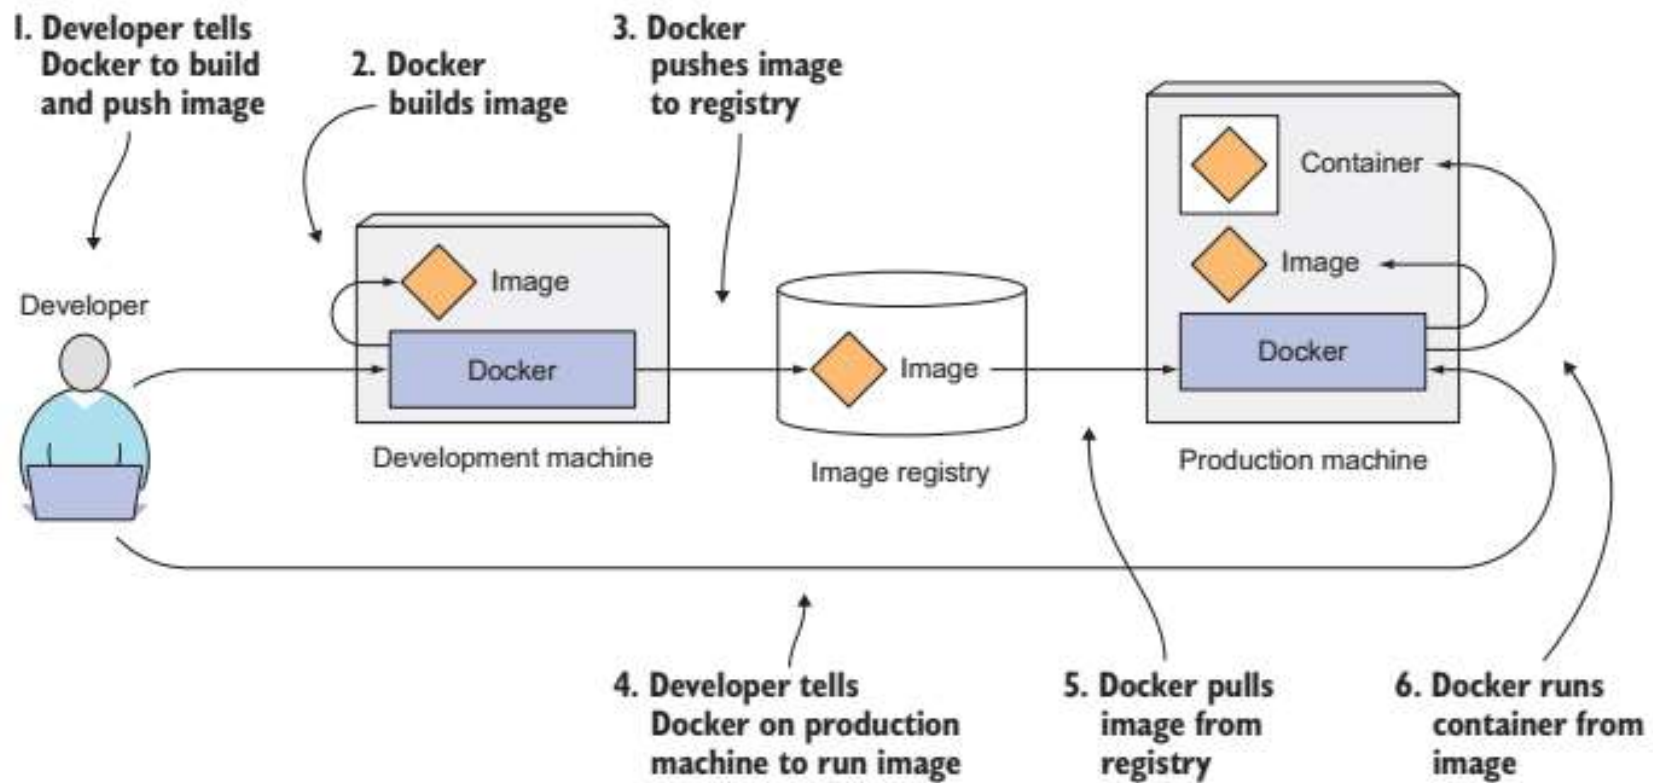


## Microservices-based application

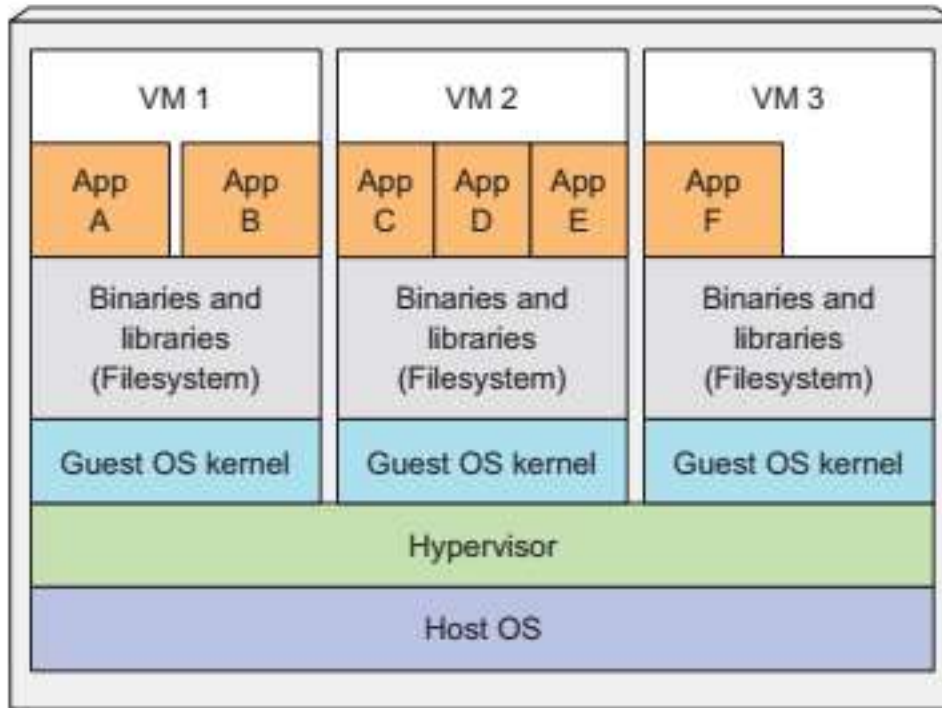


**Figure 1.1** Components inside a monolithic application vs. standalone microservices

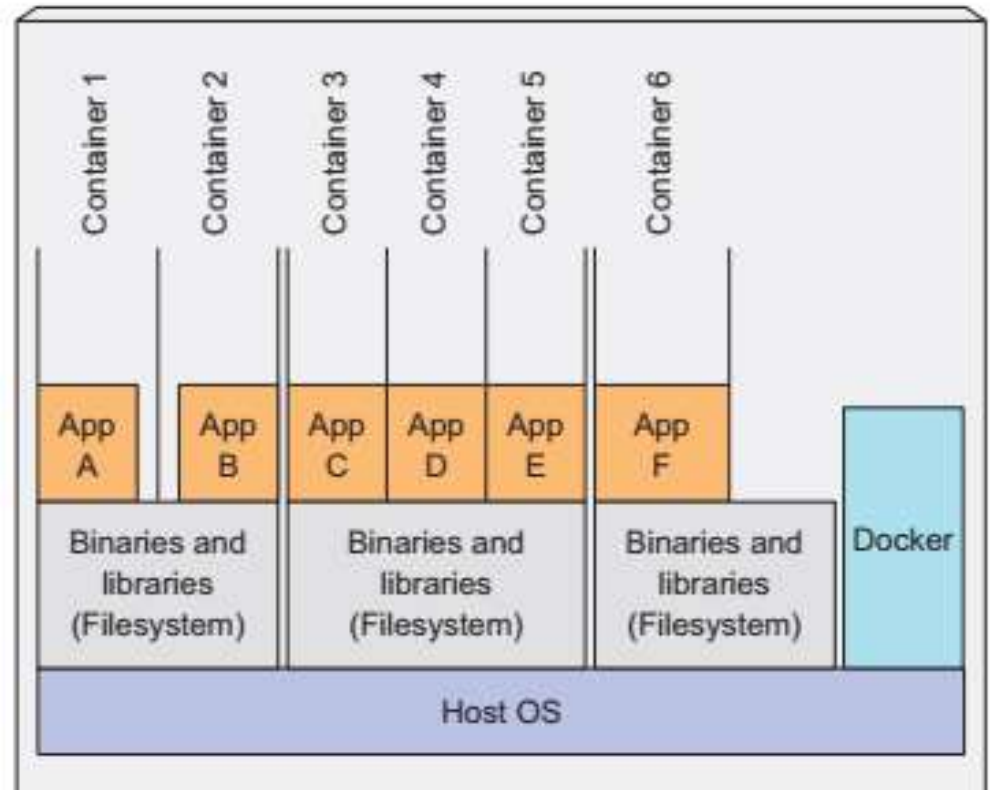


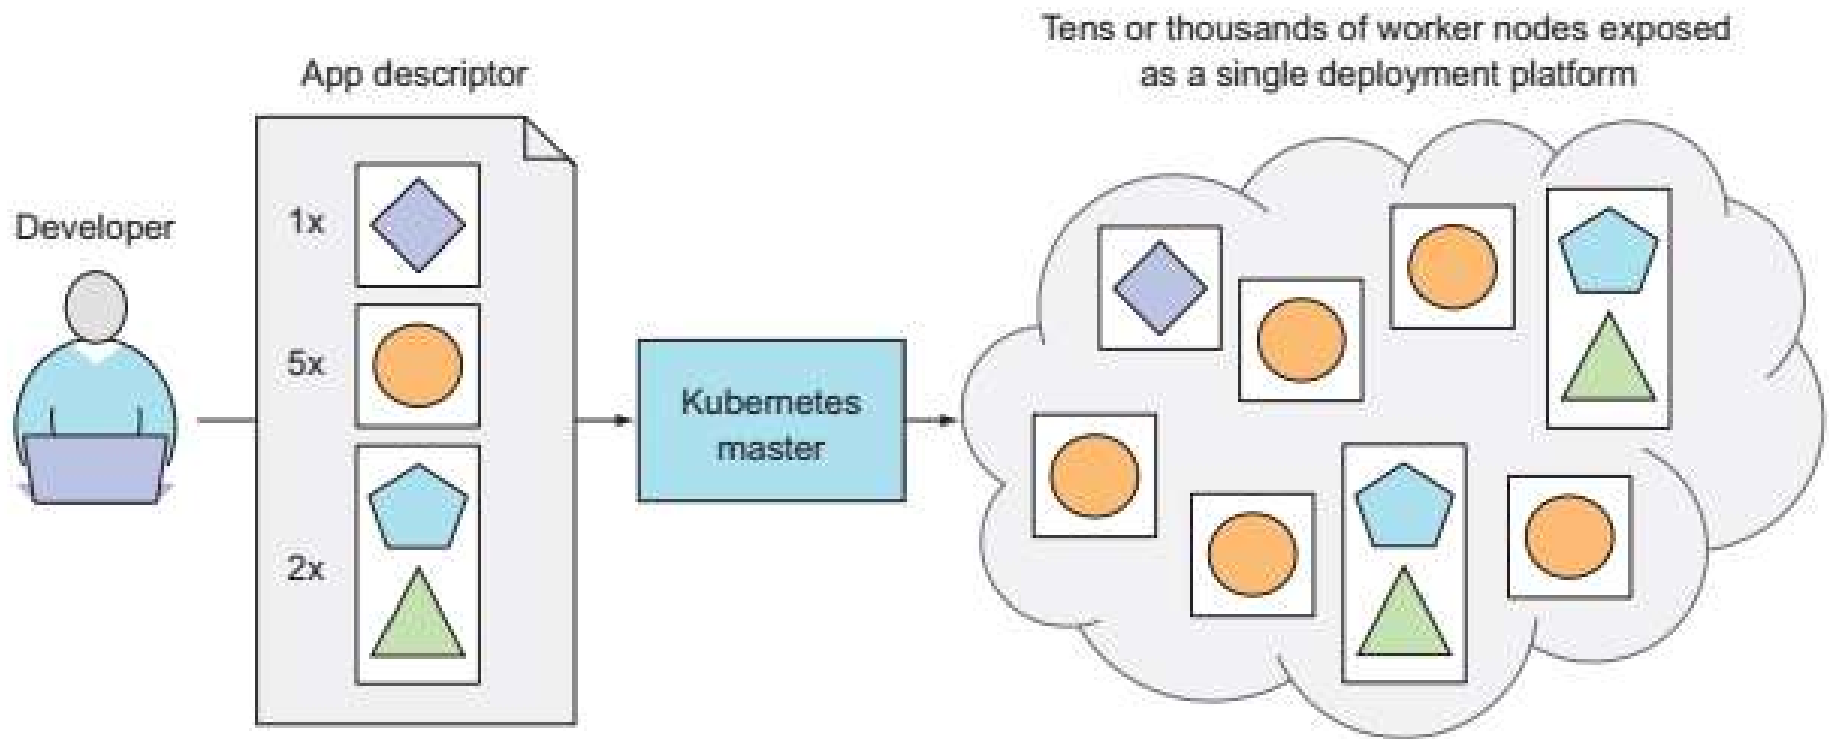


Host running multiple VMs

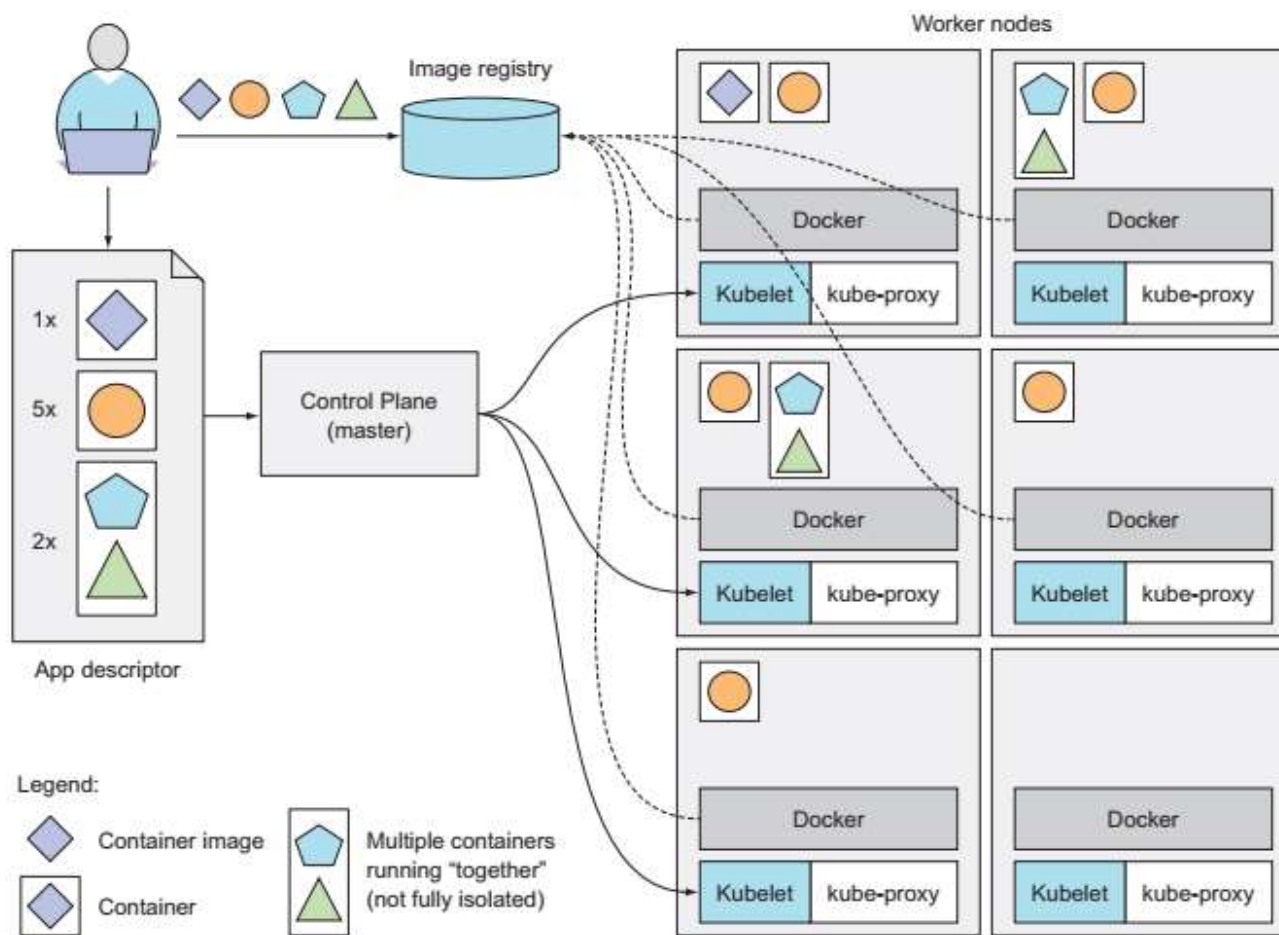


Host running multiple Docker containers

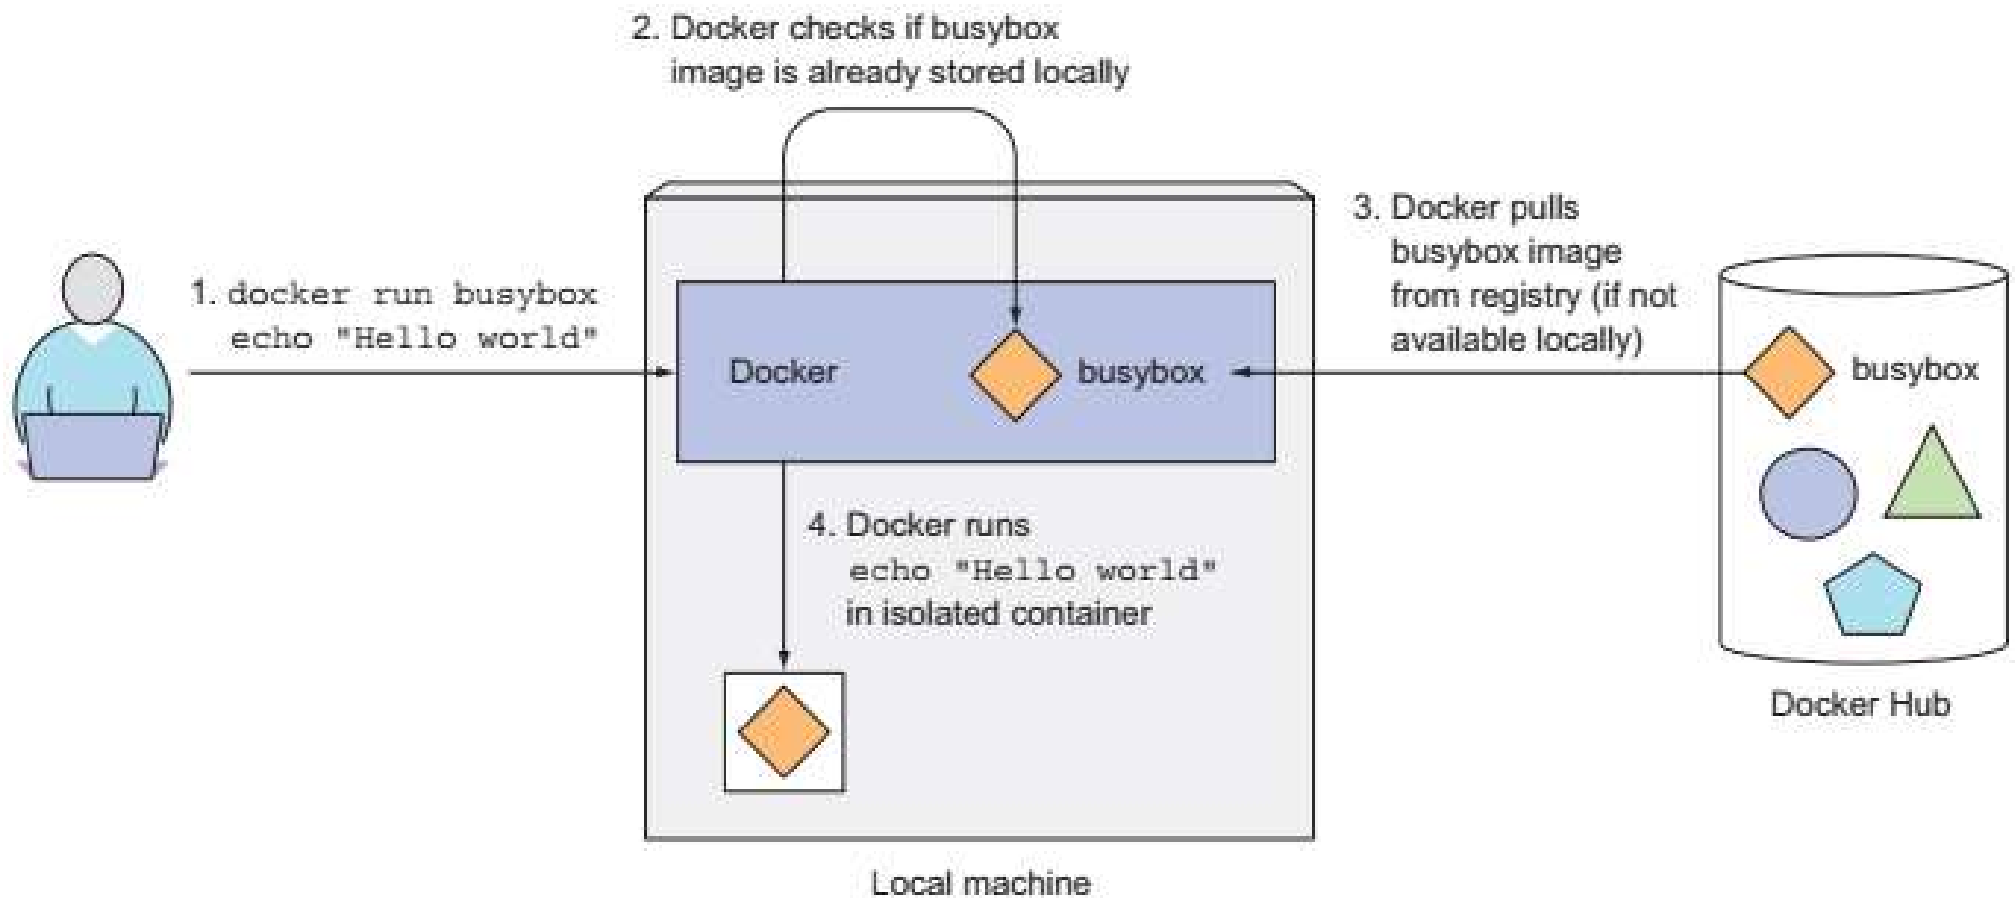




**Figure 1.8** Kubernetes exposes the whole datacenter as a single deployment platform.

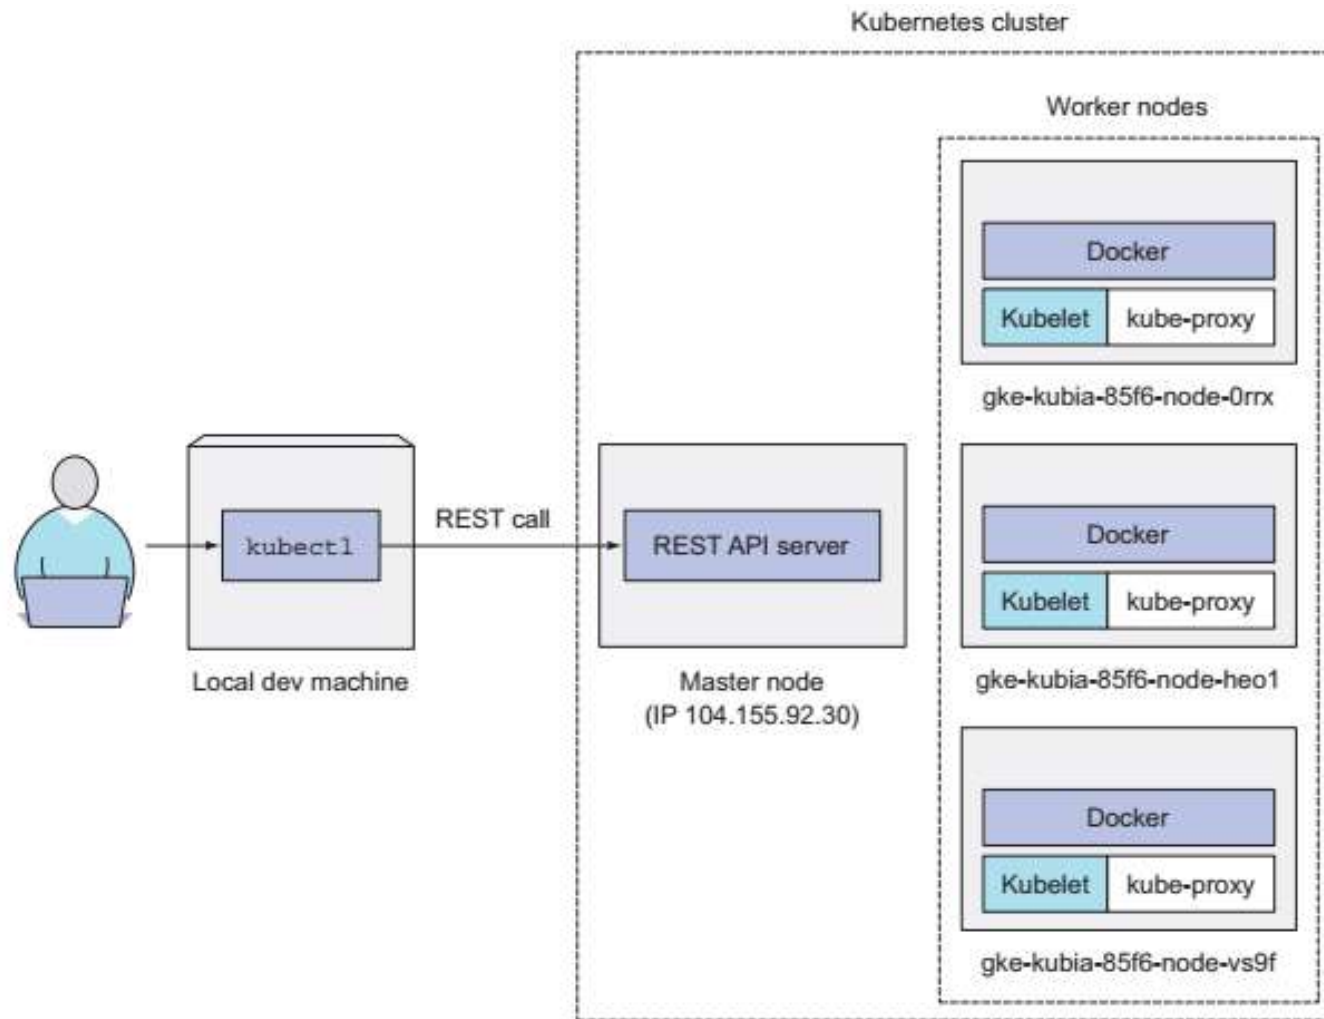


**Figure 1.10** A basic overview of the Kubernetes architecture and an application running on top of it




**Figure 2.1** Running echo “Hello world” in a container based on the busybox container image

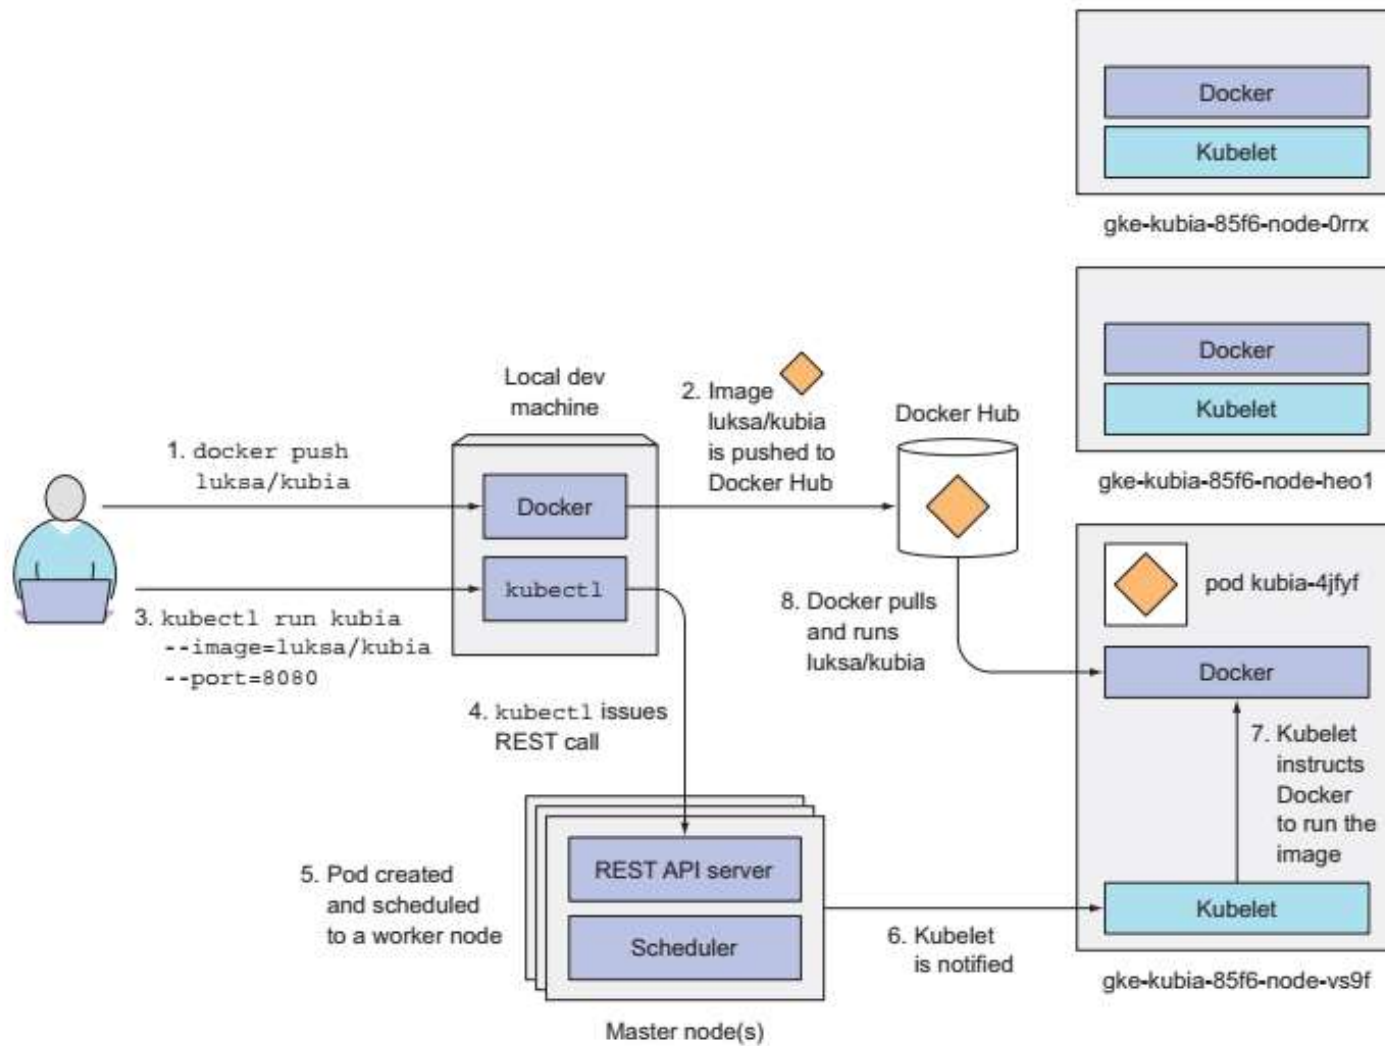





**Figure 2.4** How you're interacting with your three-node Kubernetes cluster



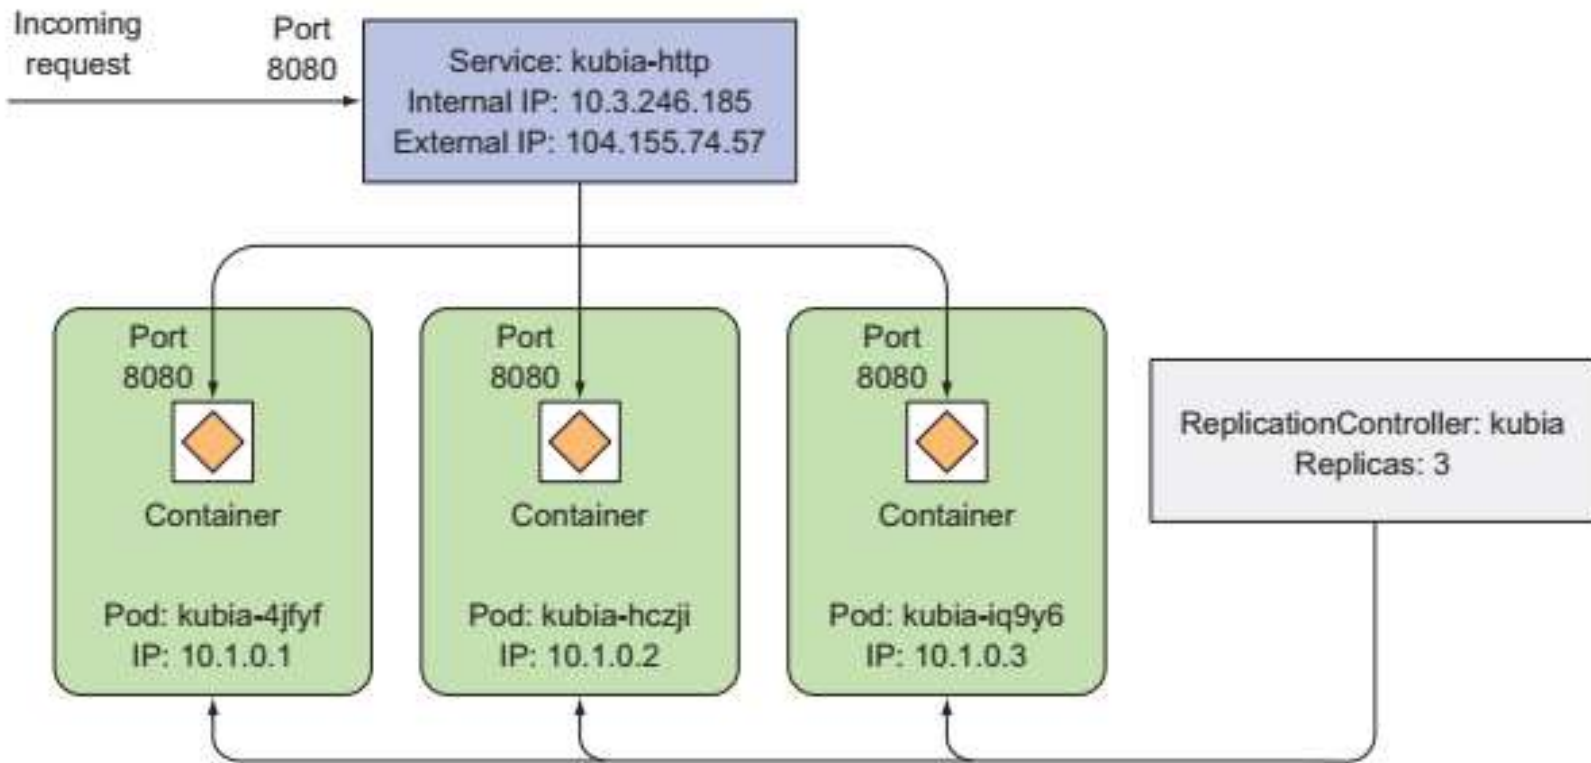
**Figure 2.5** The relationship between containers, pods, and physical worker nodes



**Figure 2.6** Running the luksa/kubia container image in Kubernetes



**Figure 2.7** Your system consists of a ReplicationController, a Pod, and a Service.



**Figure 2.8** Three instances of a pod managed by the same ReplicationController and exposed through a single service IP and port.

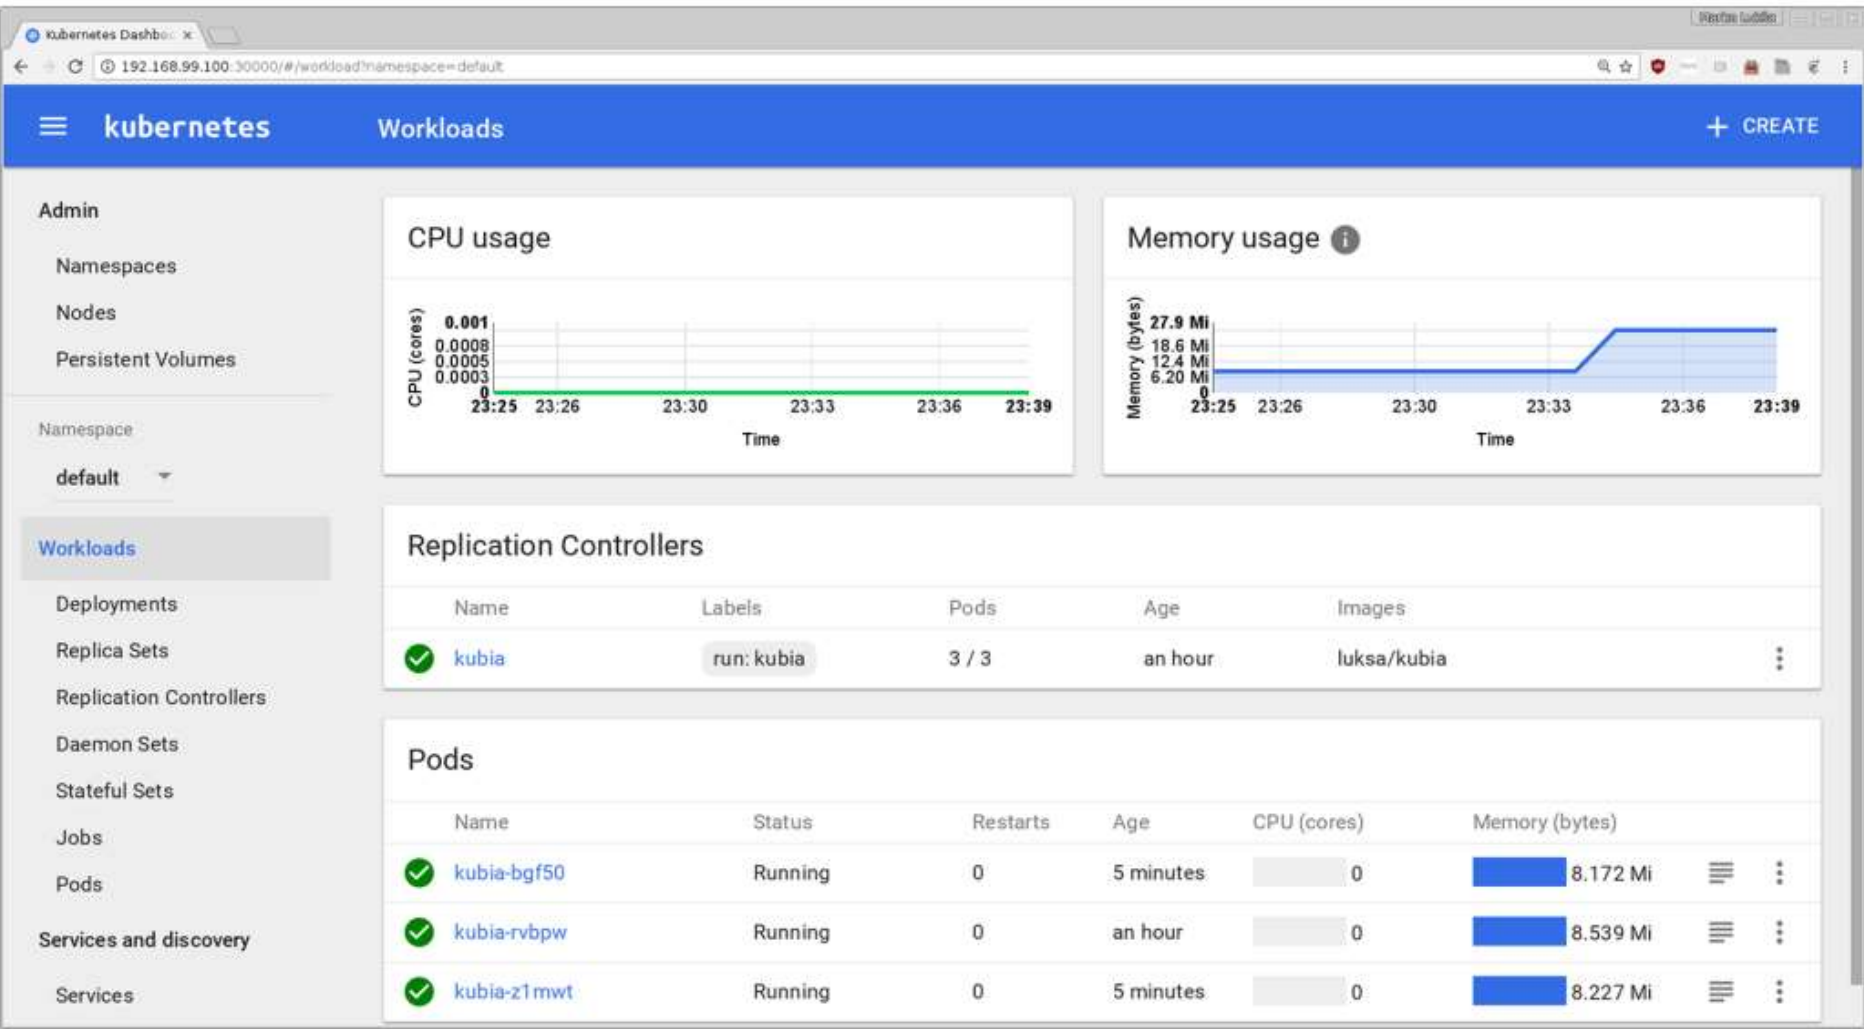
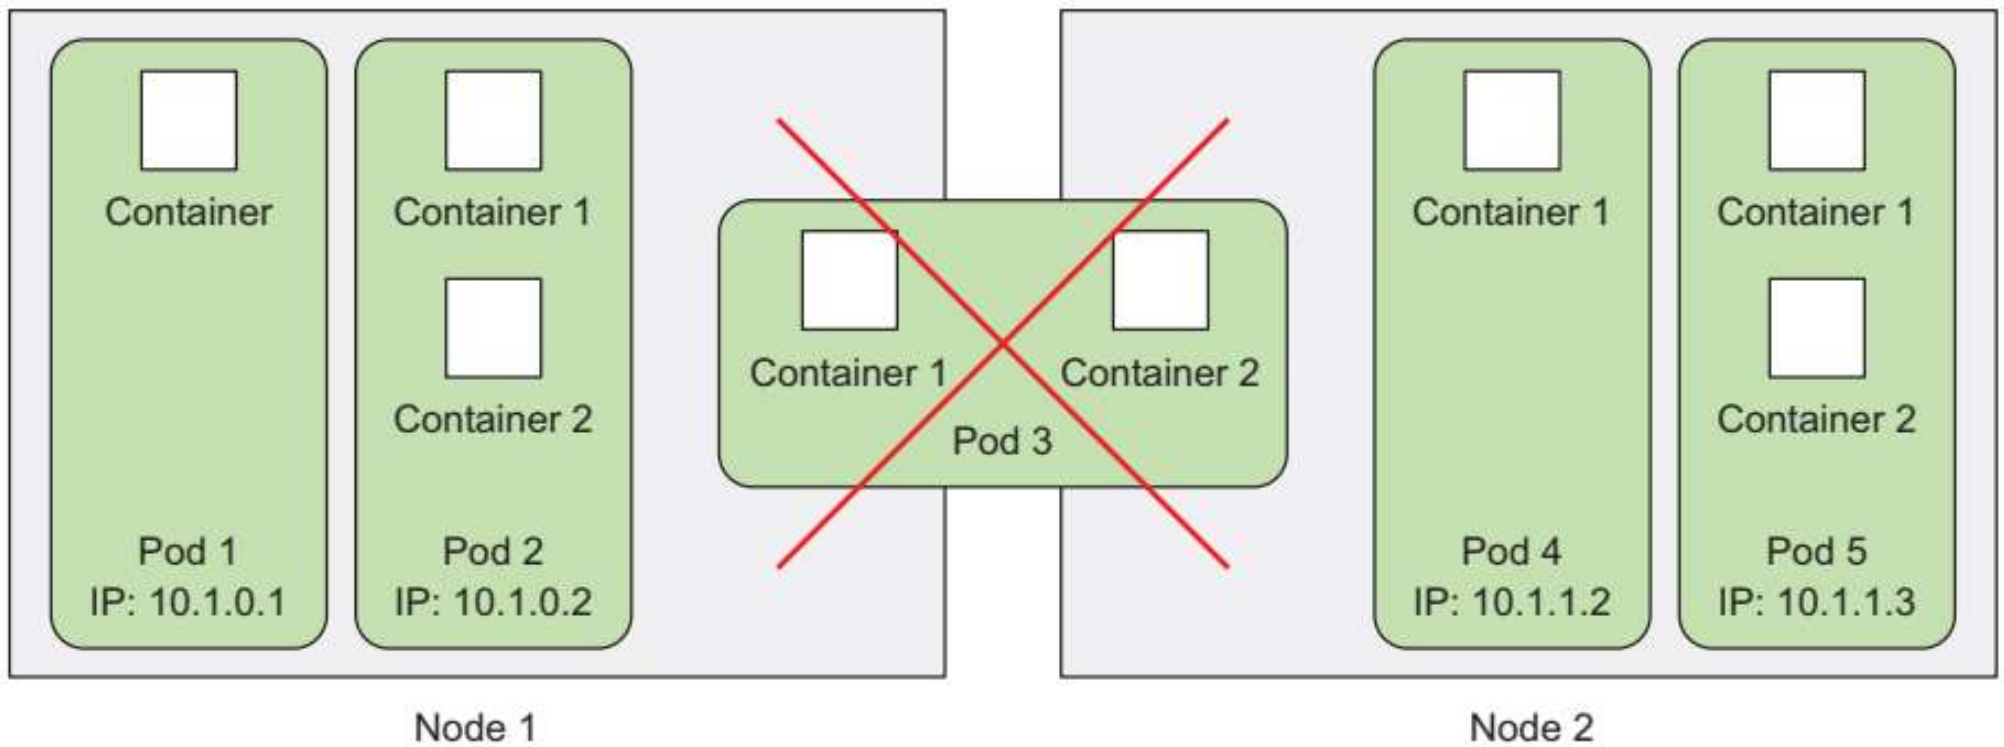
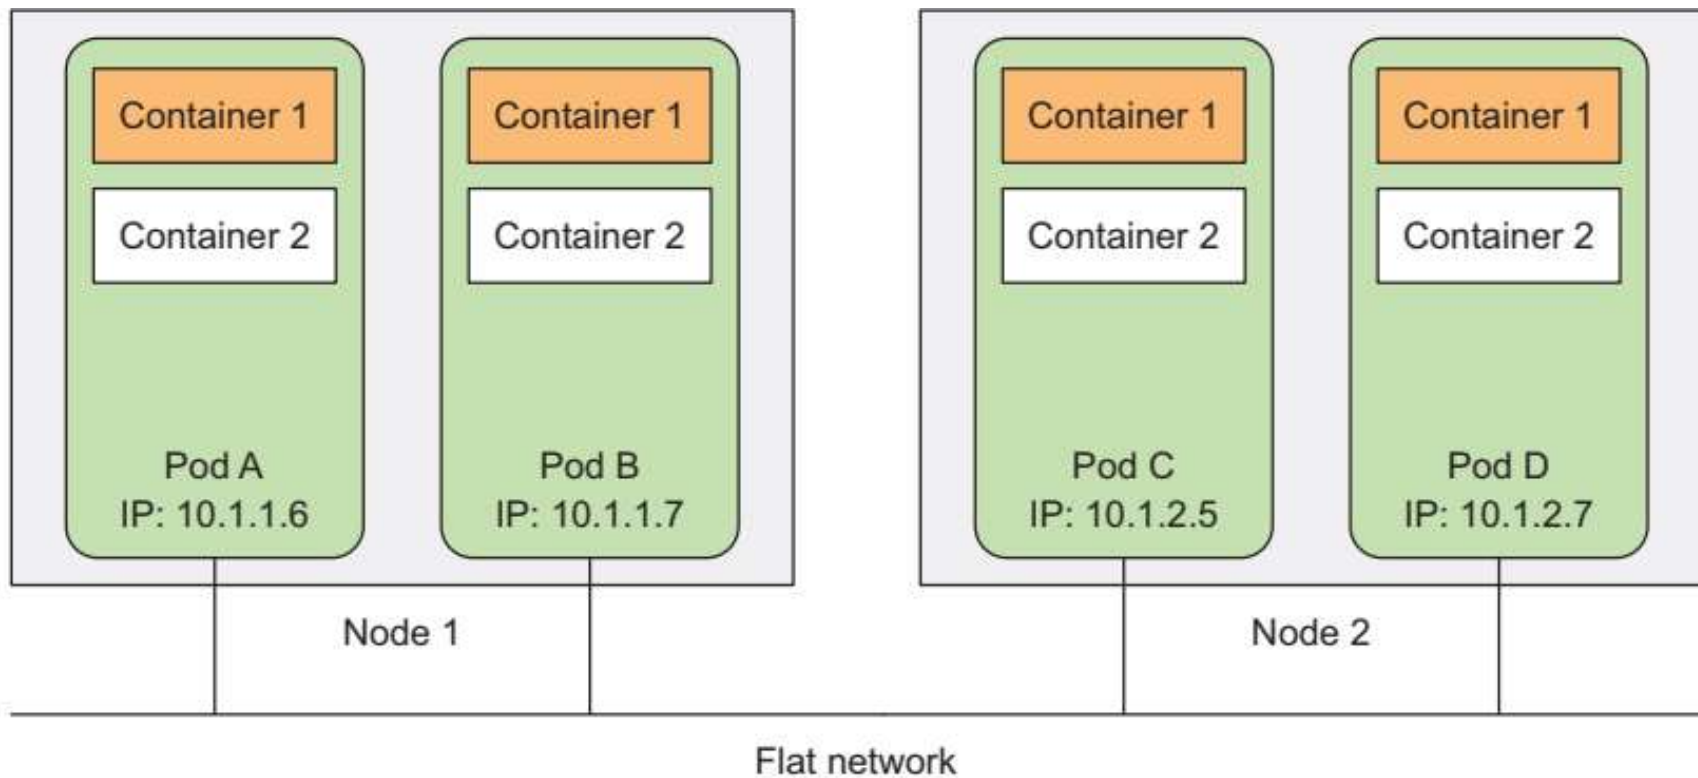


Figure 2.9 Screenshot of the Kubernetes web-based dashboard

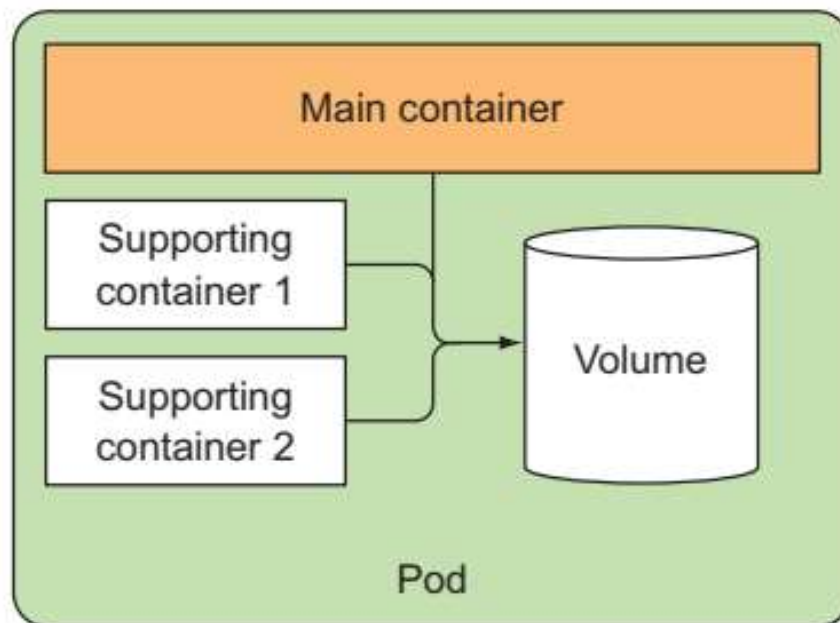


**Figure 3.1** All containers of a pod run on the same node. A pod never spans two nodes.

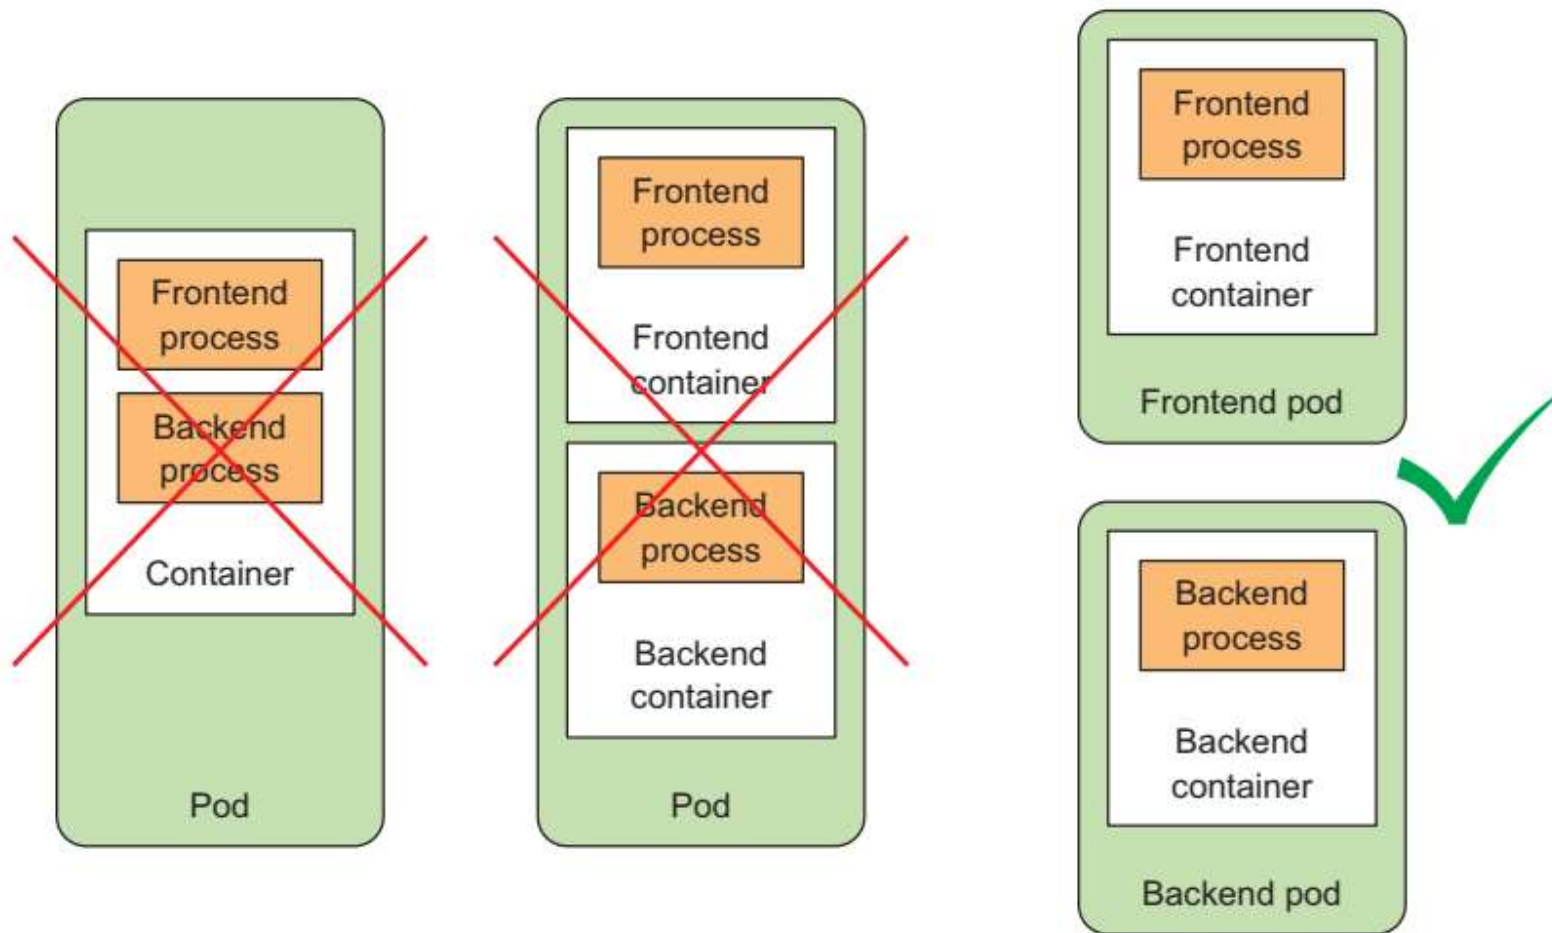


**Figure 3.2** Each pod gets a routable IP address and all other pods see the pod under that IP address.




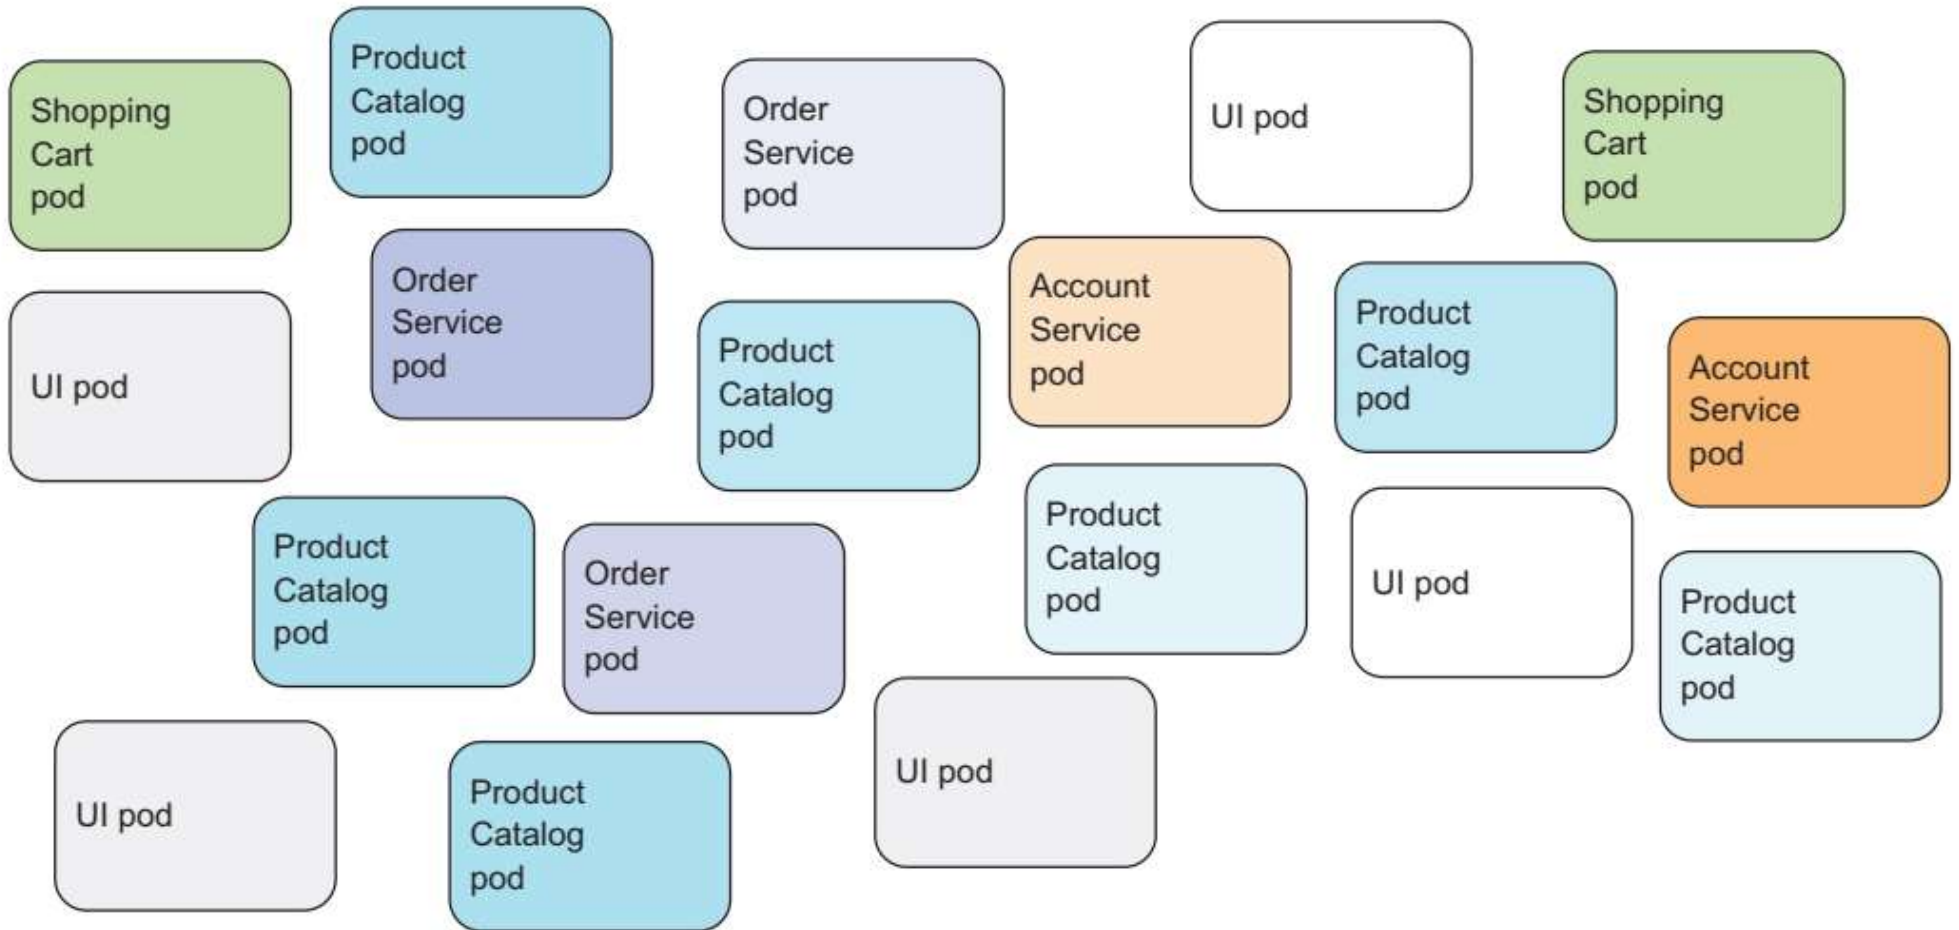


**Figure 3.3** Pods should contain tightly coupled containers, usually a main container and containers that support the main one.

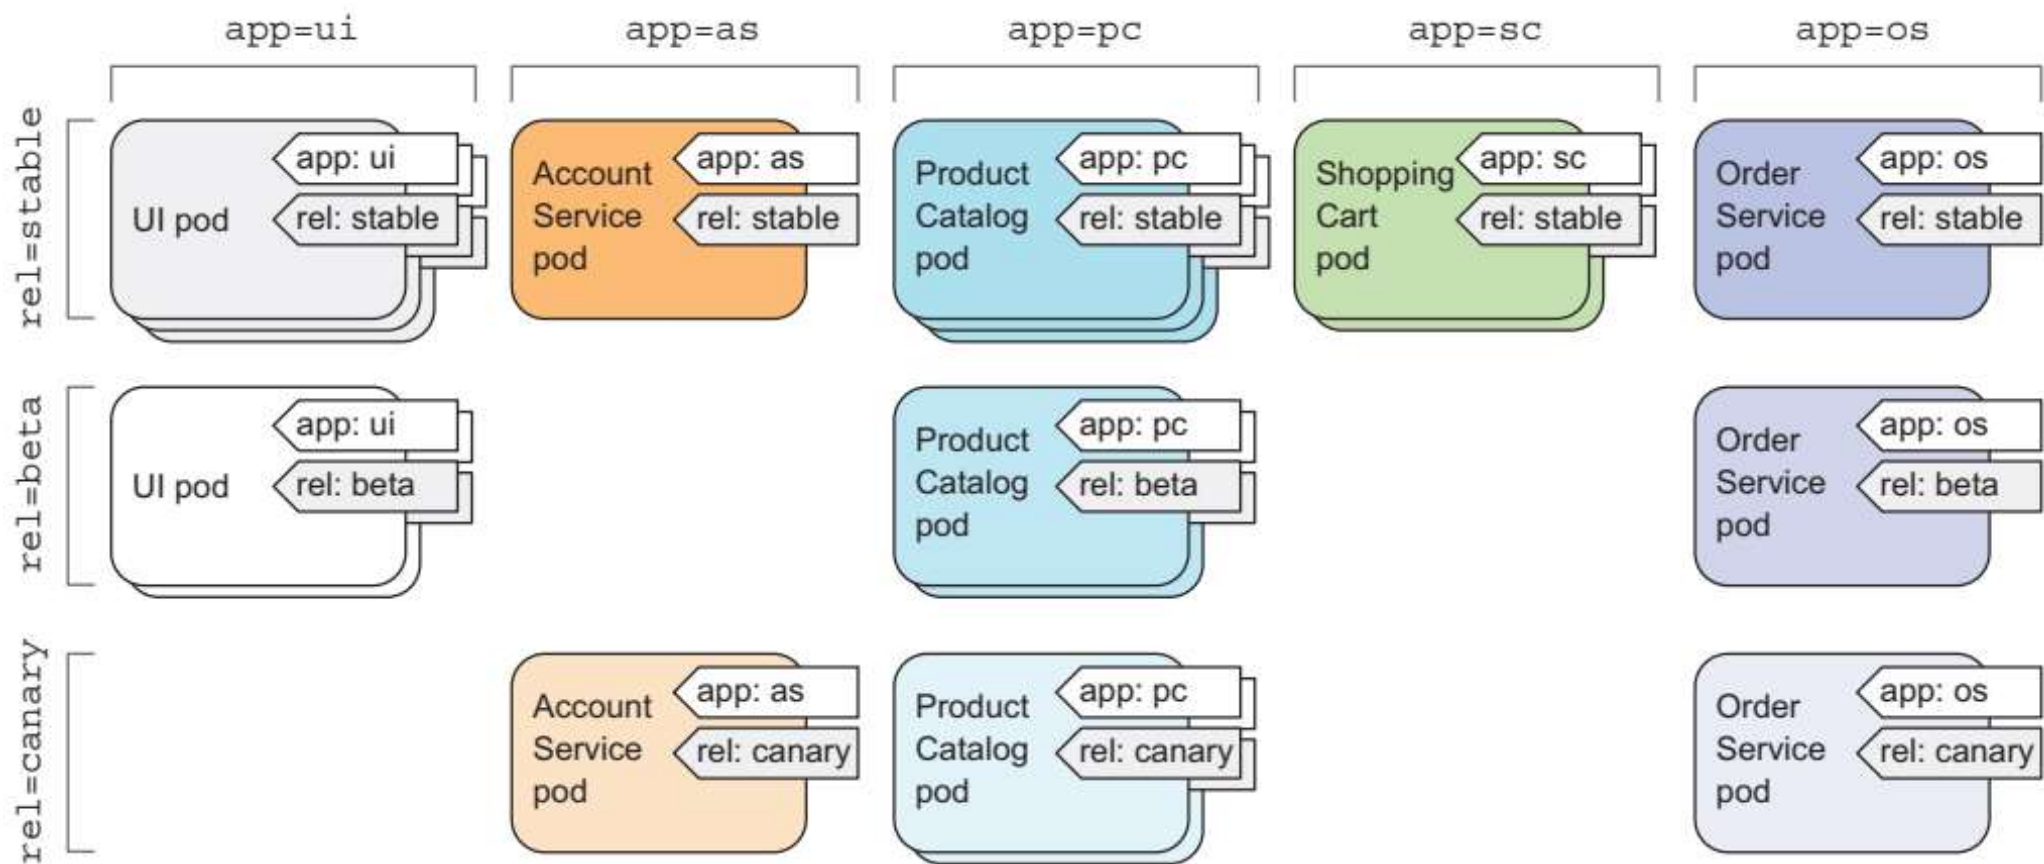


### Listing 3.2 A basic pod manifest: kuba-manual.yaml

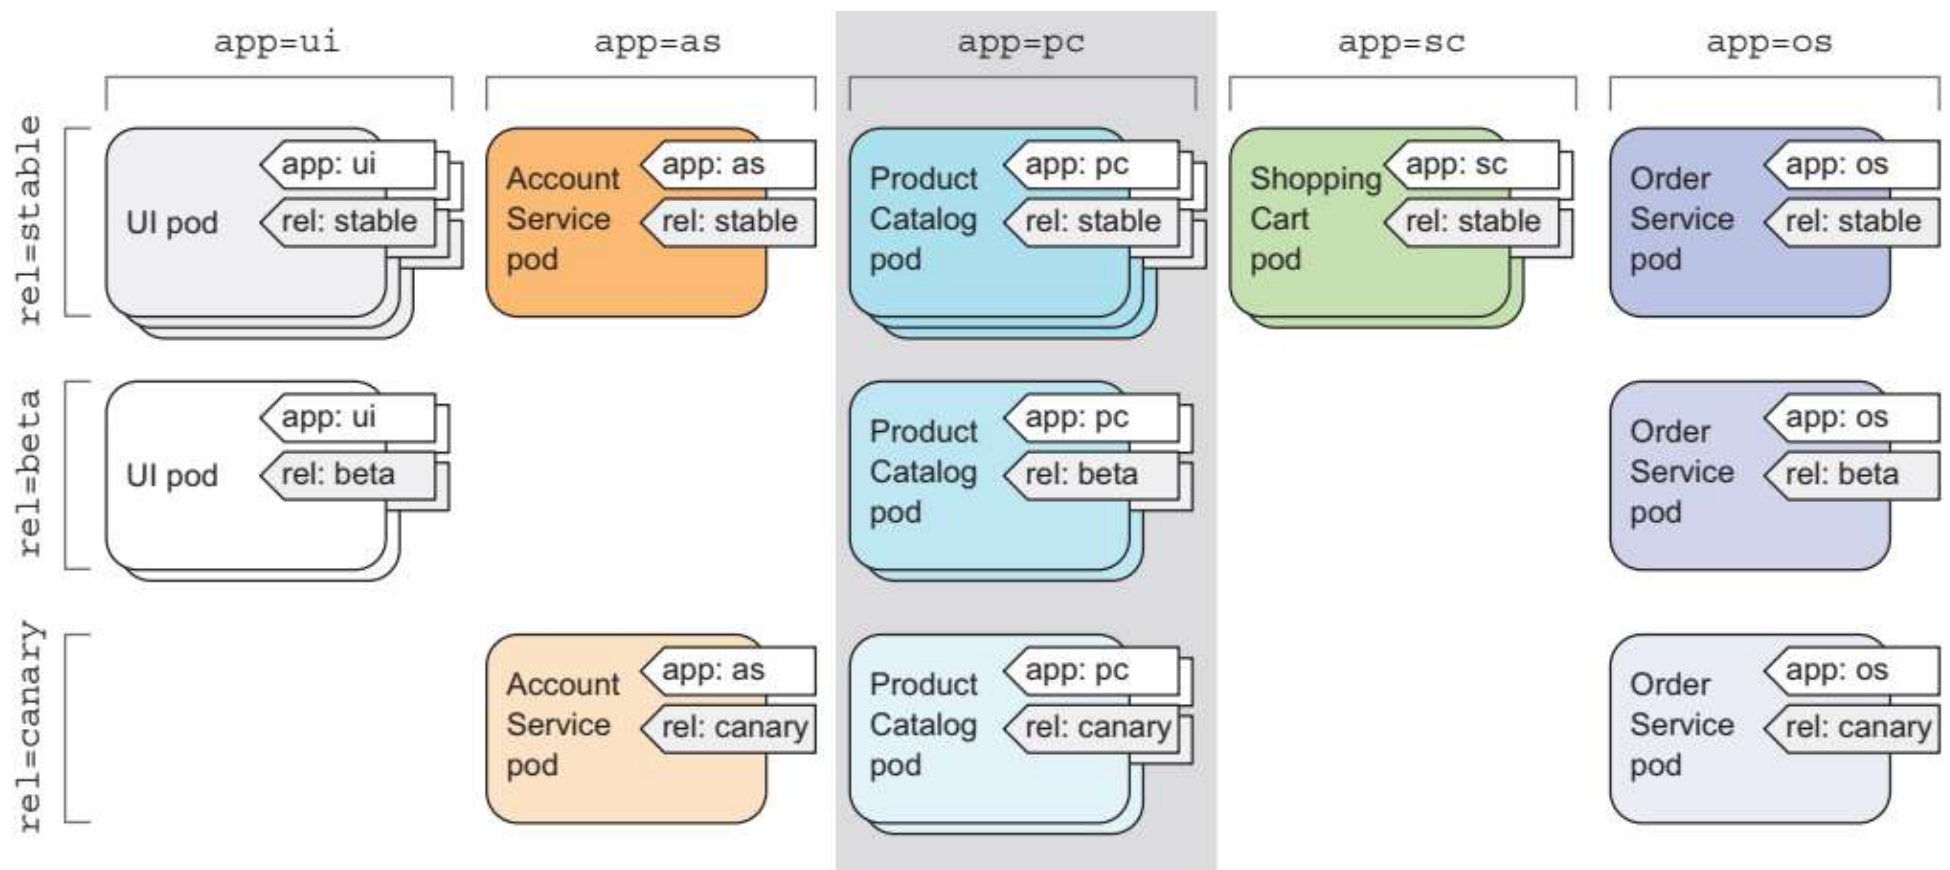




**Figure 3.6** Uncategorized pods in a microservices architecture

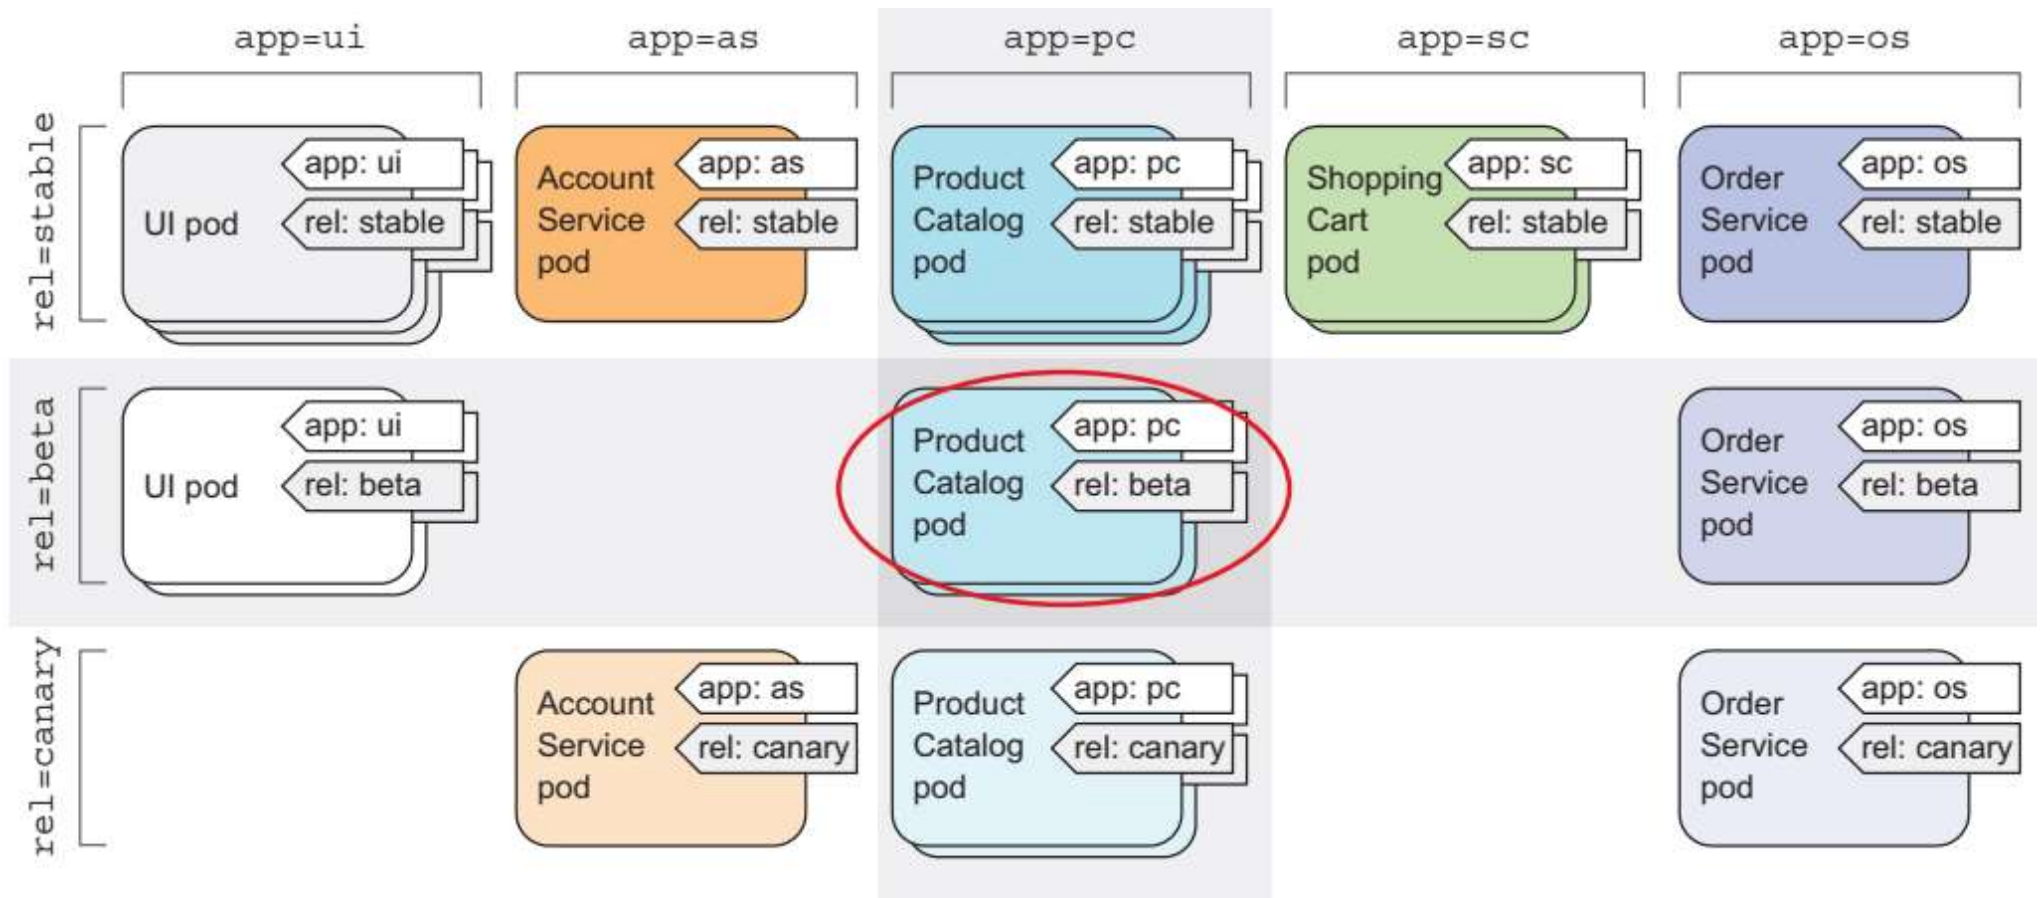


**Figure 3.7** Organizing pods in a microservices architecture with pod labels

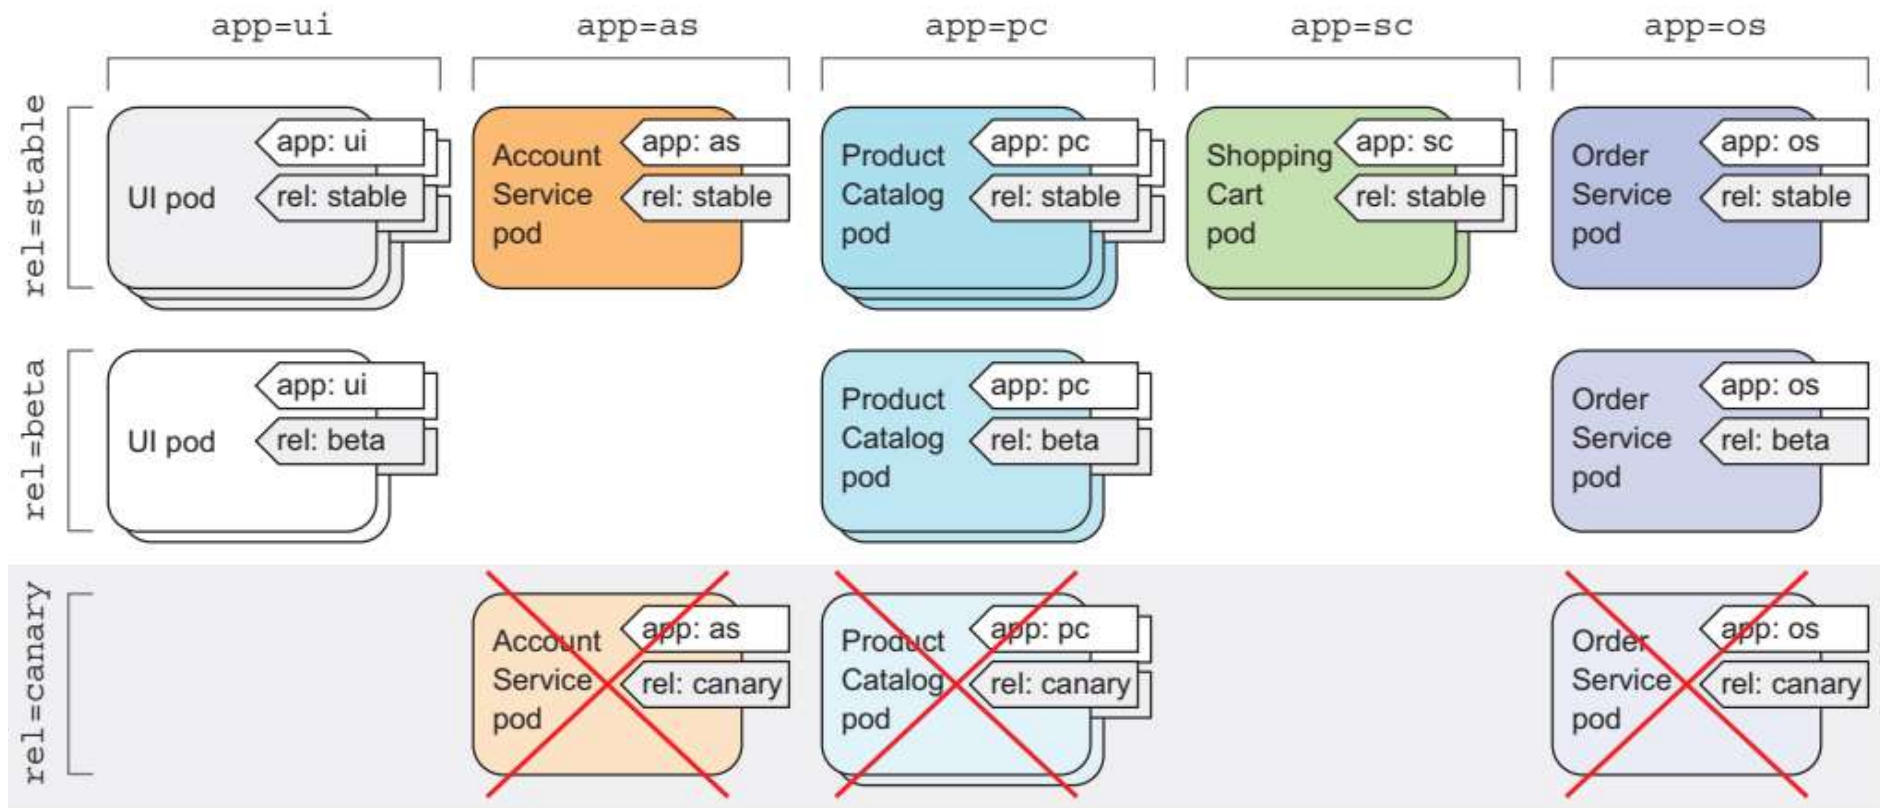


**Figure 3.8** Selecting the product catalog microservice pods using the “app=pc” label selector



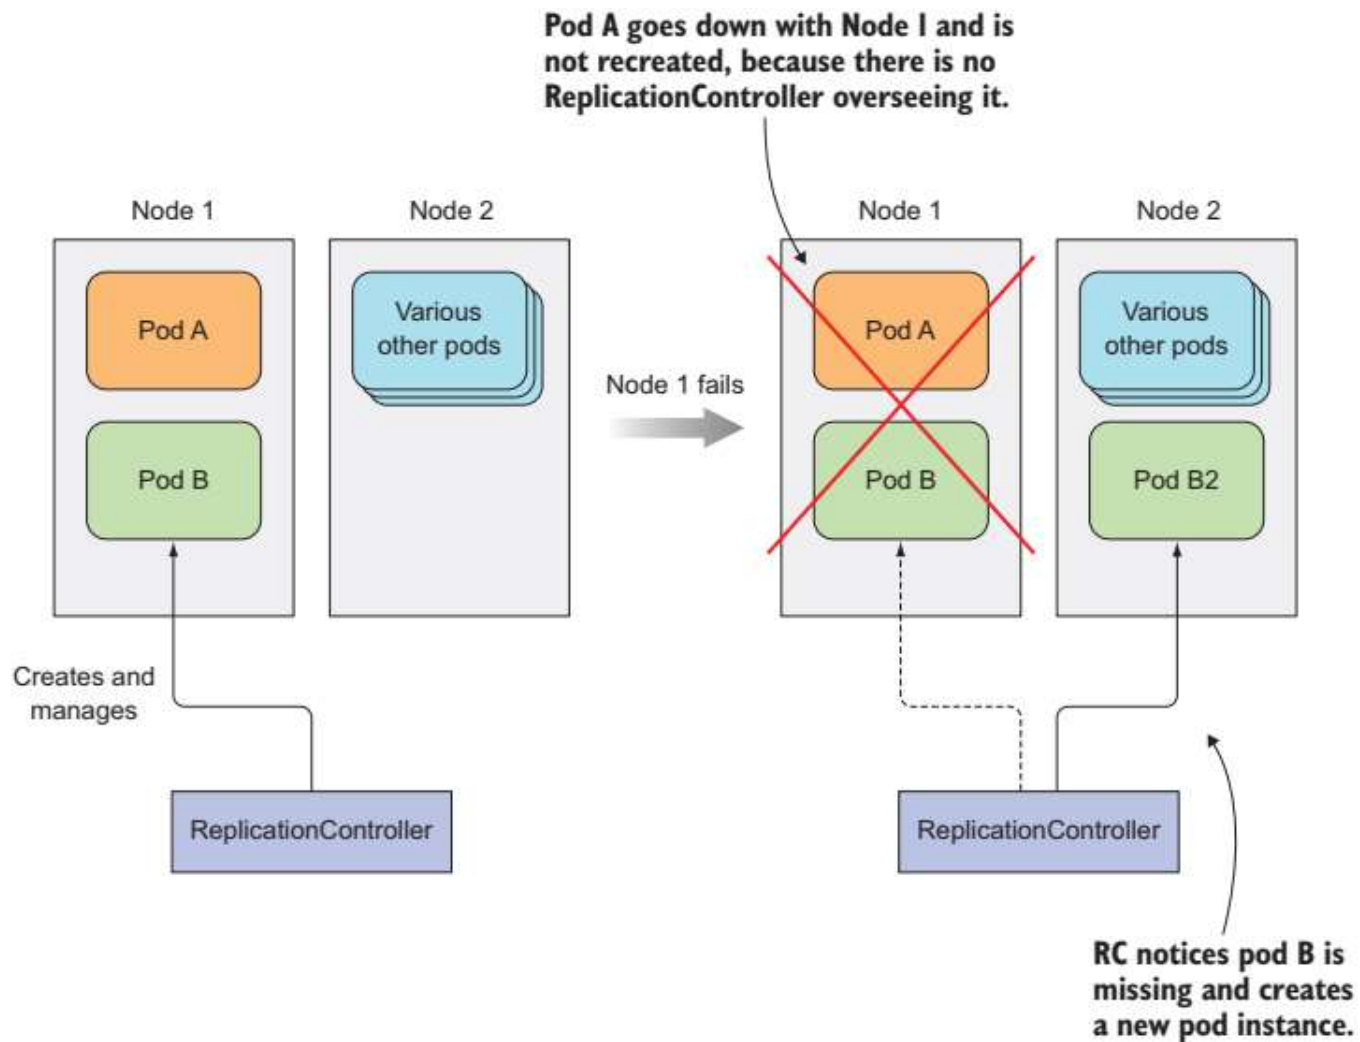


**Figure 3.9** Selecting pods with multiple label selectors

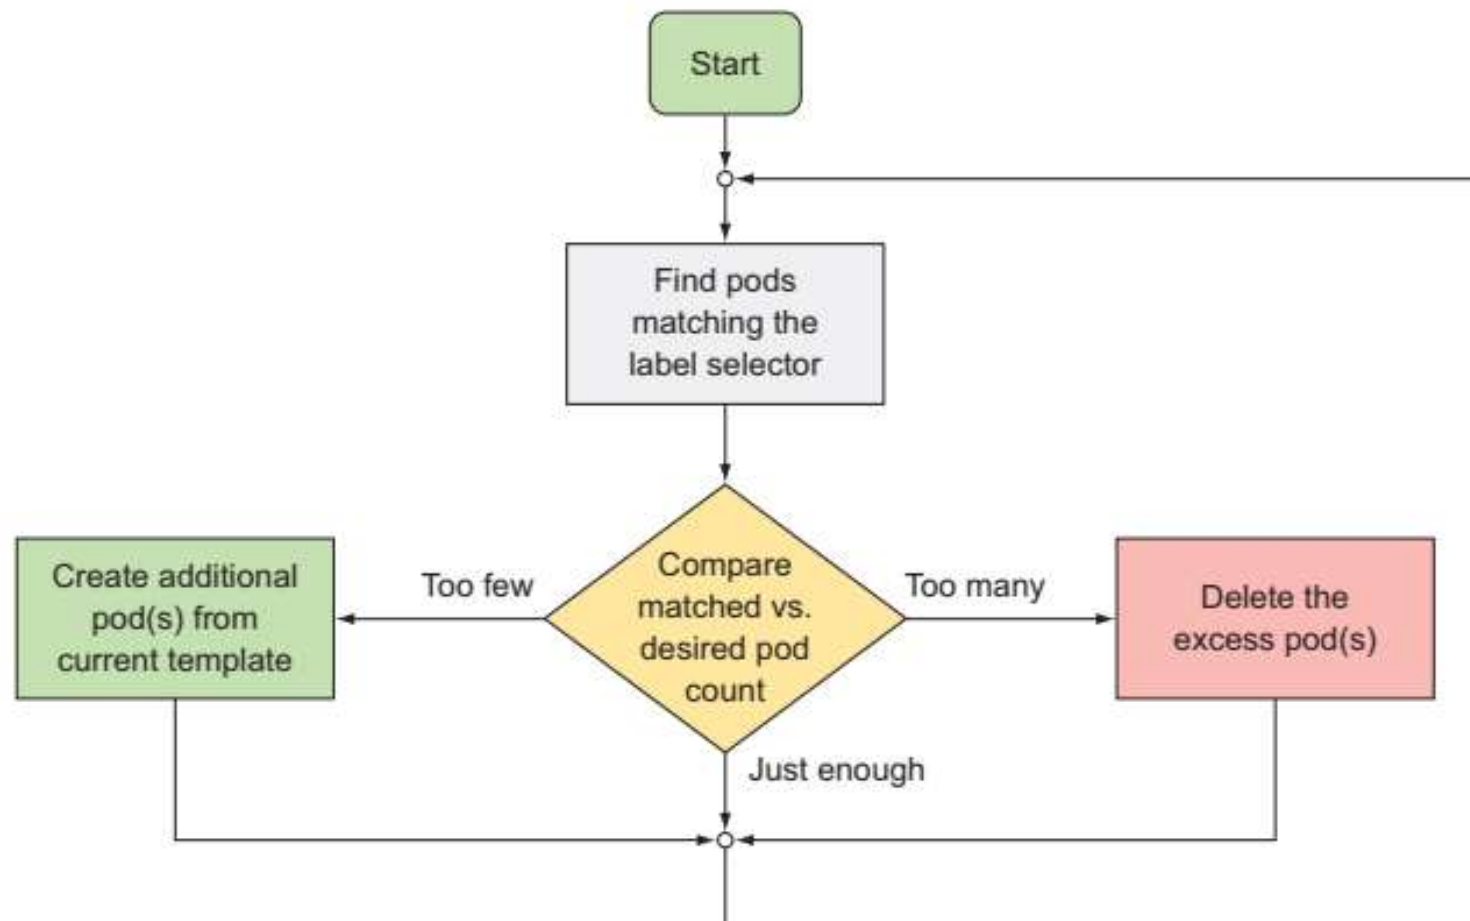


**Figure 3.10** Selecting and deleting all canary pods through the `rel=canary` label selector

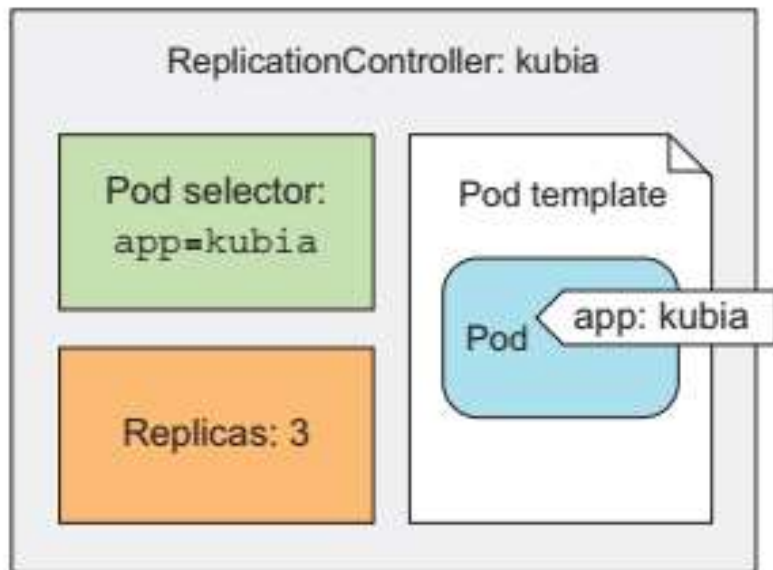




**Figure 4.1** When a node fails, only pods backed by a ReplicationController are recreated.



**Figure 4.2 A ReplicationController's reconciliation loop**



**Figure 4.3** The three key parts of a `ReplicationController` (pod selector, replica count, and pod template)

#### Listing 4.4 A YAML definition of a ReplicationController: kubia-rc.yaml

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: kubia
spec:
  replicas: 3
  selector:
    app: kubia
```

This manifest defines a  
ReplicationController (RC)

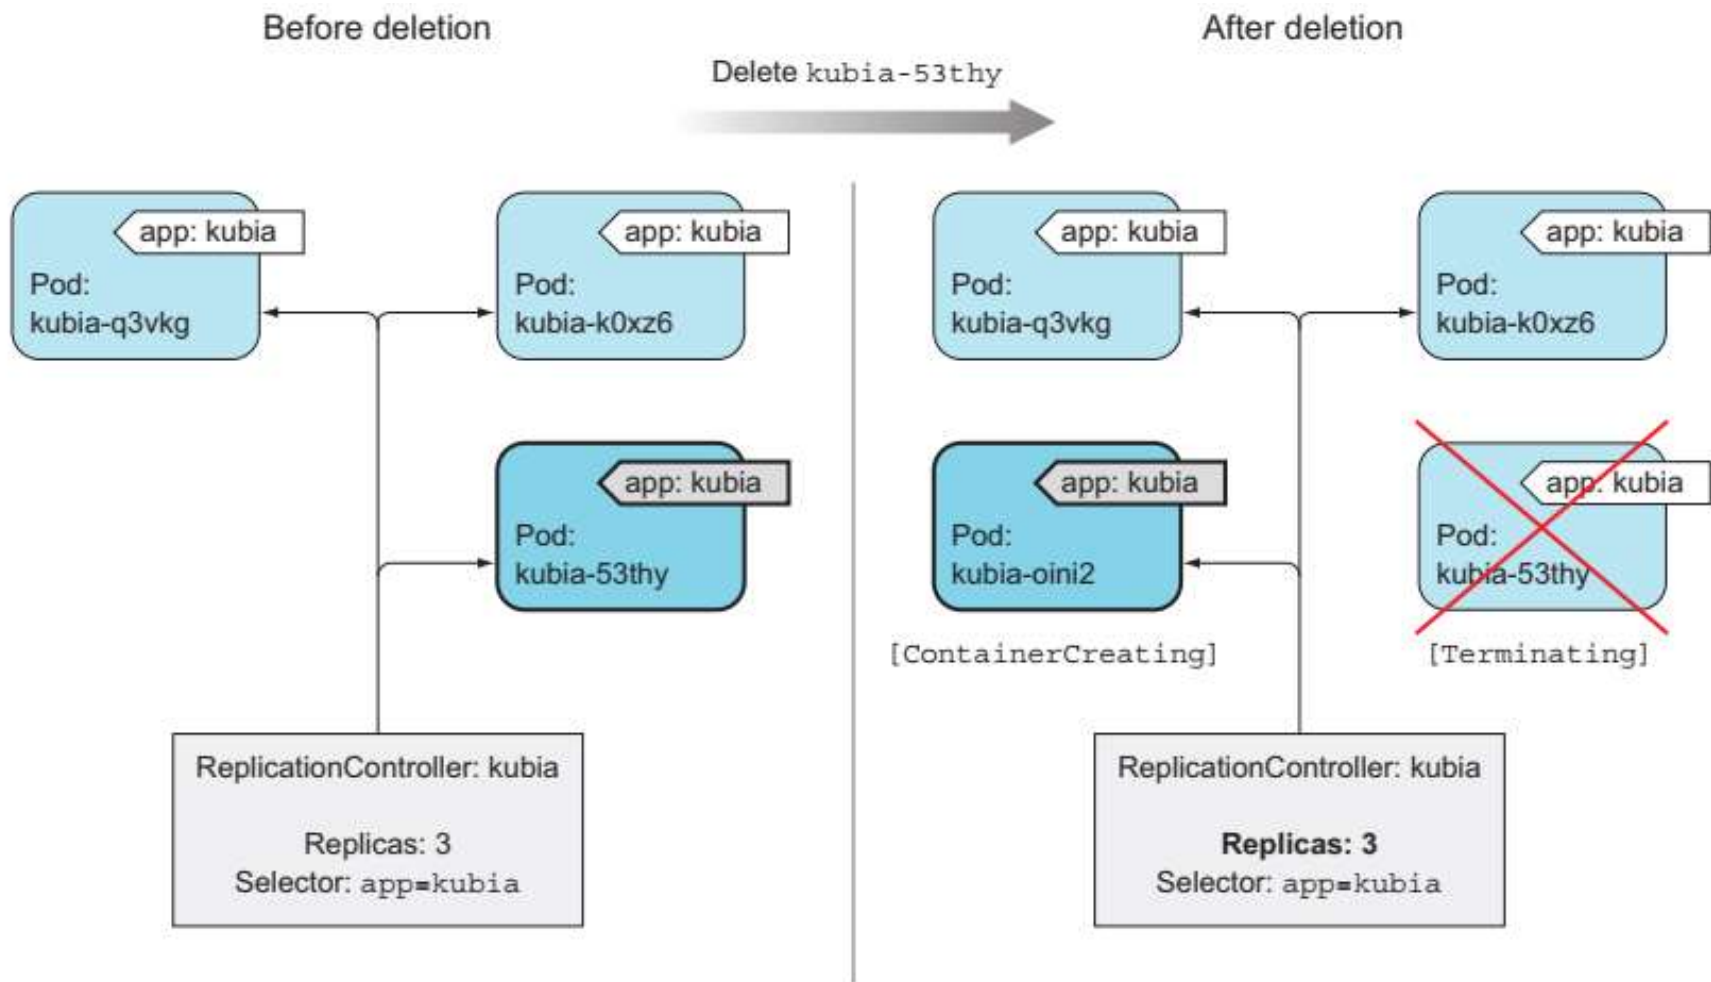
The name of this  
ReplicationController

The desired number  
of pod instances

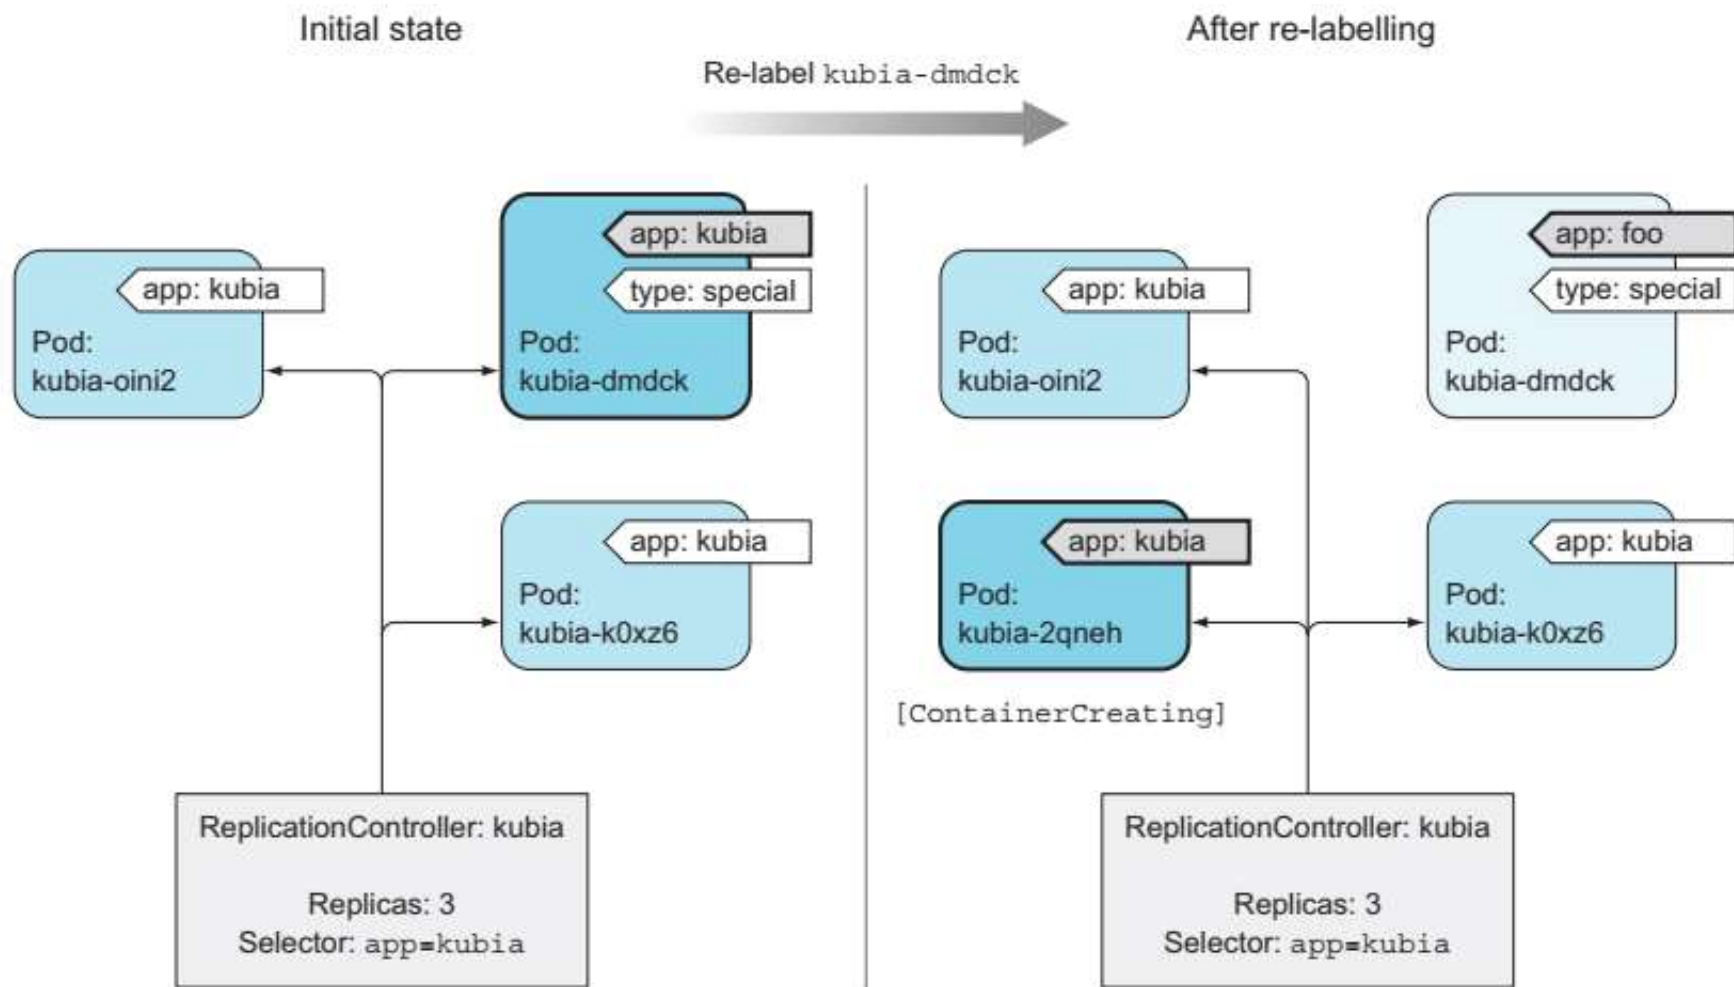
The pod selector determining  
what pods the RC is operating on

```
template:
  metadata:
    labels:
      app: kubia
  spec:
    containers:
      - name: kubia
        image: luksa/kubia
        ports:
          - containerPort: 8080
```

The pod template  
for creating new  
pods

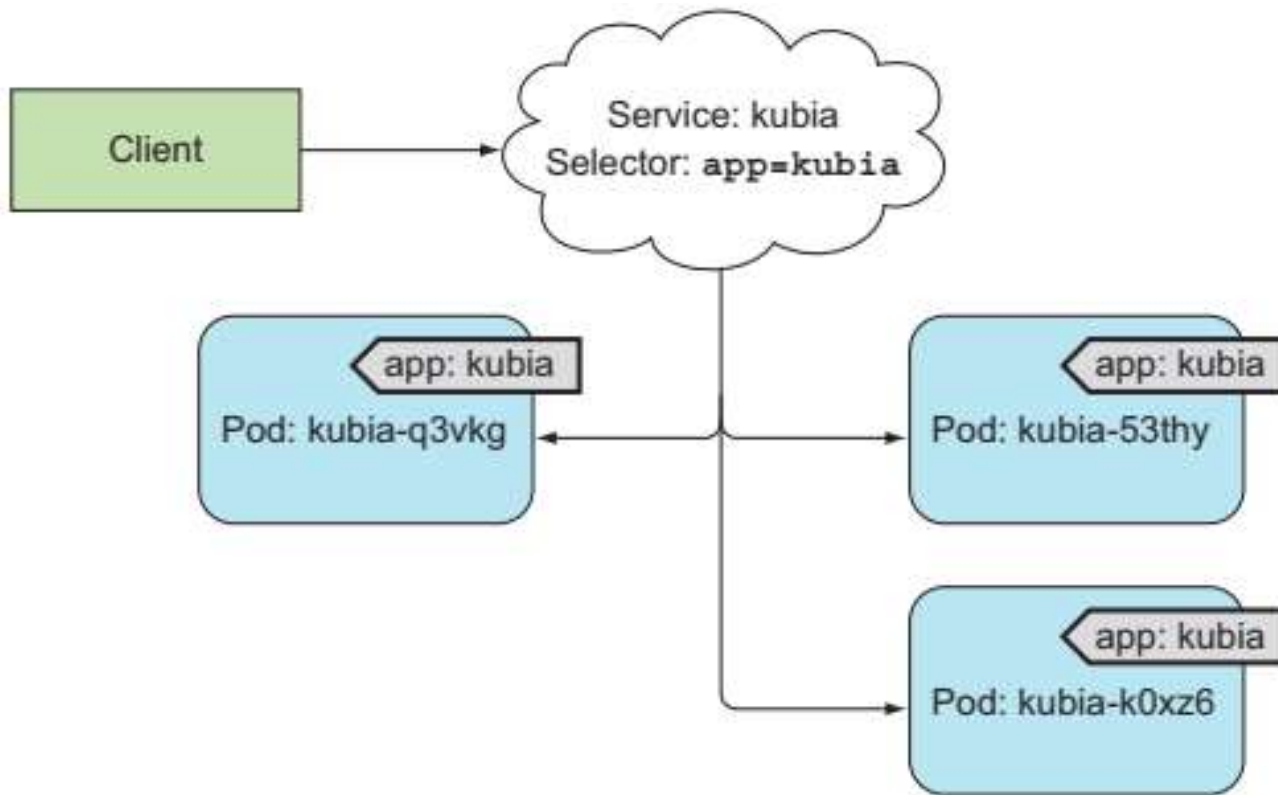


**Figure 4.4** If a pod disappears, the ReplicationController sees too few pods and creates a new replacement pod.



**Figure 4.5** Removing a pod from the scope of a ReplicationController by changing its labels

# Services



**Figure 5.2** Label selectors determine which pods belong to the Service.




### Listing 5.1 A definition of a service: kubia-svc.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: kubia
spec:
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: kubia
```

**The port this service  
will be available on**

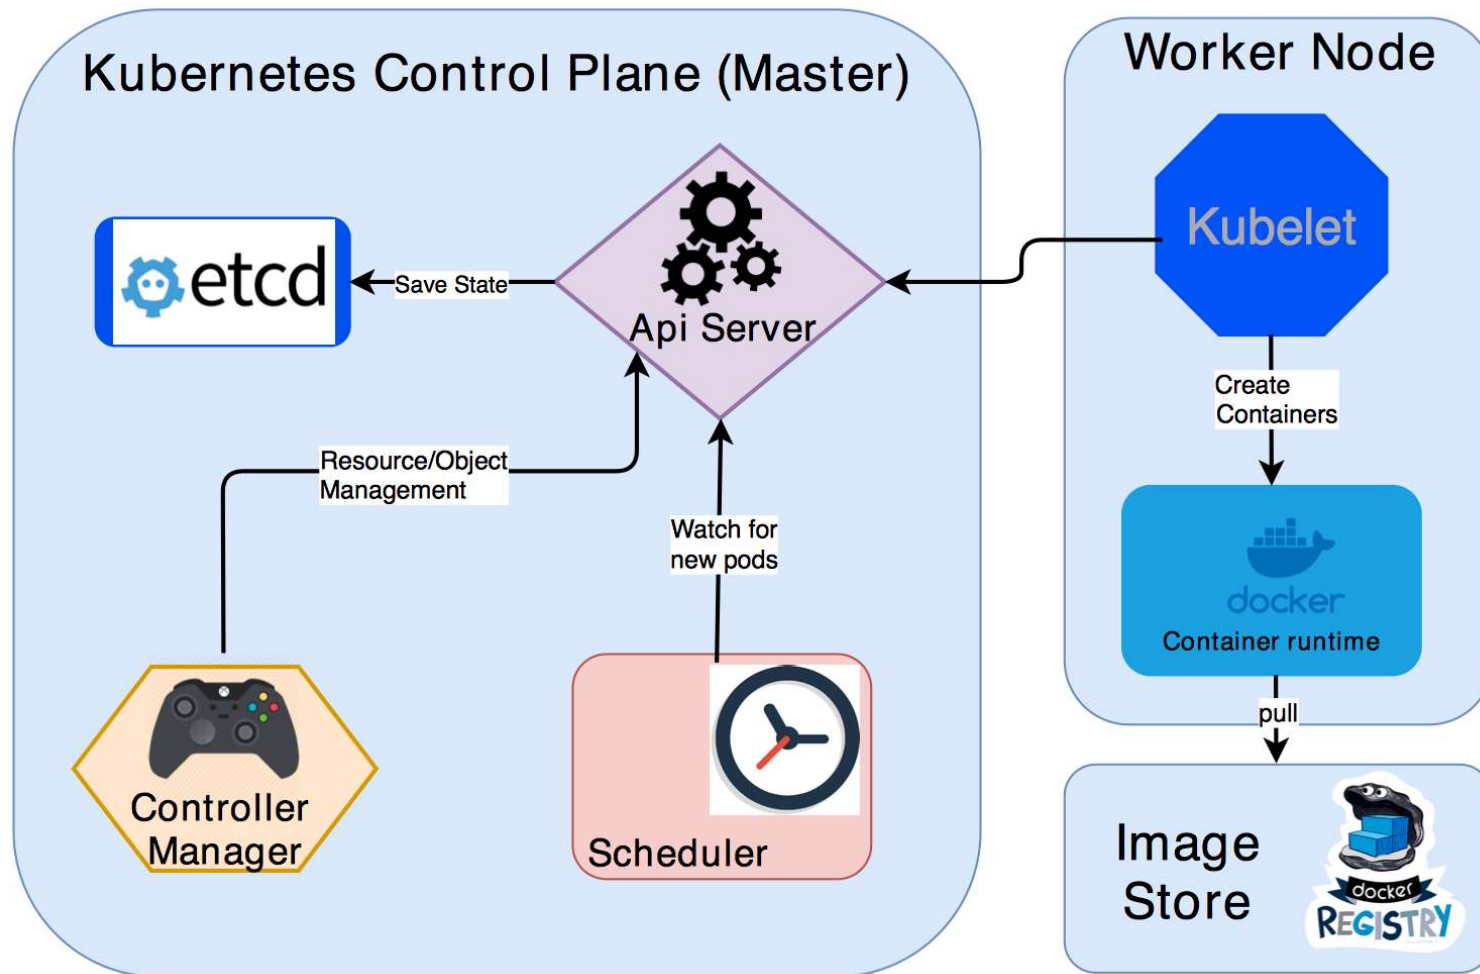
**The container port the  
service will forward to**

**All pods with the app=kubia  
label will be part of this service.**



**Figure 5.5** Exposing a service to external clients

# Kubernetes Architecture



# Kubernetes Architecture

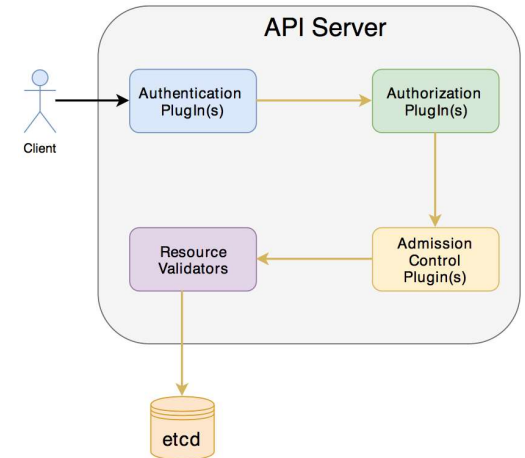
- A Kubernetes cluster consists of two main components:
  - Master (Control Plane)
  - Worker Nodes.
- Master has following components. These components are responsible for maintaining the state of the cluster:
  - etcd distributed key value store.
  - API Server.
  - Controller Manager
  - Scheduler
- Every worker node consists of the following components.
- These components are responsible for deploying and running the application containers.
  - Kubelet
  - Container Runtime (Docker)

# Master Components

- etcd:
  - Stores the cluster status and metadata.
  - A distributed key value store
  - Provides reliable way of storing data across a cluster of machines.
  - API Server directly talk to etcd store.
  - K8s stores all its data under /registry directory in etcd.
- Api Server:
  - The central place for all other components.
  - Api Server will take care about validating the object before saving the information to etcd.
  - The client for the Api Server can be either kubectl (command line tool) or a Rest Api client.

# Master Components

- Api Server:
  - Several plugin's are invoked by Api Server before creating/deleting/updating the object in etcd.
- Scheduler:
  - Allocate what node the pods needs to be created
- Controller Manager:
  - Make sure the actual state of the system converges towards the desired state.
  - Watch the API Server for changes to resources/objects and perform necessary actions like create/update/delete of the resource.



# Worker Node components

- Kubelet
  - The agent that runs on each node in the cluster.
  - Monitors the Api Server for Pods
  - Start the pod's containers by instructing to docker runtime.
  - Monitors the status of running containers and reports to api server
  - Also do health checks for the container and restart if needed.
- Docker
  - Container runtime used by Kubelet for spinning up Containers

# Other components

- Nodes:
  - Machine on which Kubernetes is installed.
  - This is where containers inside the pods will be launched by Kubernetes.
- Master Node:
  - Responsible for managing the cluster
  - A Kubernetes cluster also contains one or more master nodes that run the Kubernetes control plane
- Pod
  - Smallest deployable unit that can be managed by Kubernetes.
  - A logical group of one or more containers that share the same IP address

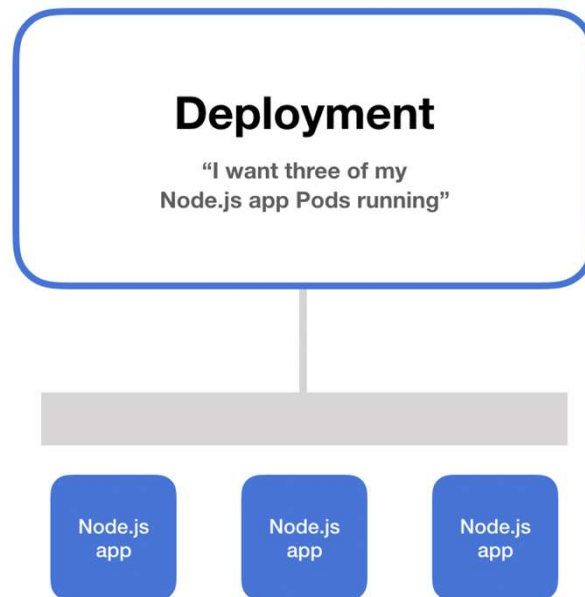


# Kubernetes namespaces

- Provide for a scope of Kubernetes resource, carving up your cluster in smaller units
  - `$ kubectl get ns`
  - `$ kubectl describe ns default`
  - `$ kubectl create namespace test`

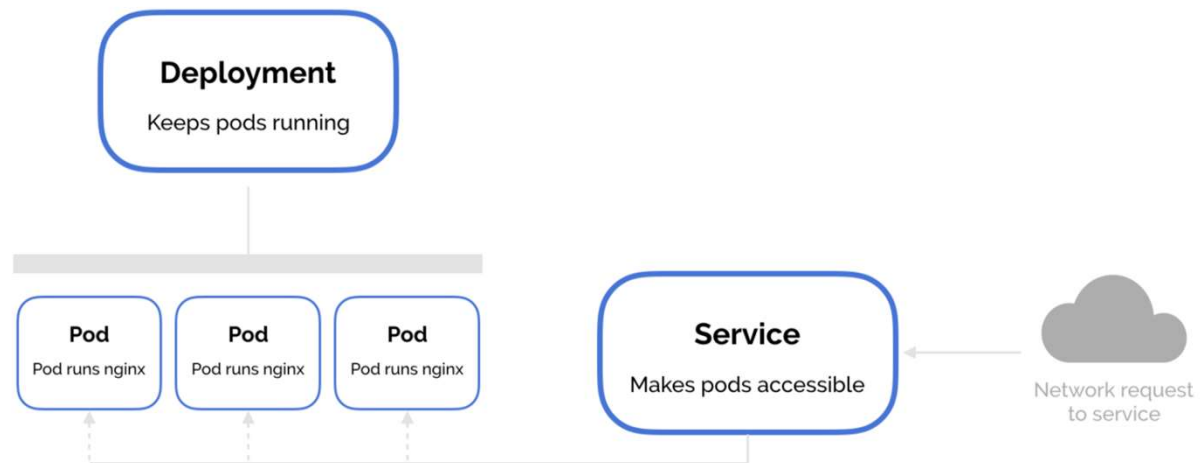
# Kubernetes deployment

- Everyone running applications on Kubernetes cluster uses a deployment.
- It's what you use to scale, roll out, and roll back versions of your applications.
- With a deployment, you tell Kubernetes how many copies of a Pod you want running. The deployment takes care of everything else.



# Deployment vs service

- A deployment is used to keep a set of pods running by creating pods from a template.
- A service is used to allow network access to a set of pods.
- To access a Deployment with one or many PODs, you need a Kubernetes Service endpoint mapped to the deployment using labels and selectors.



# Installation - Online

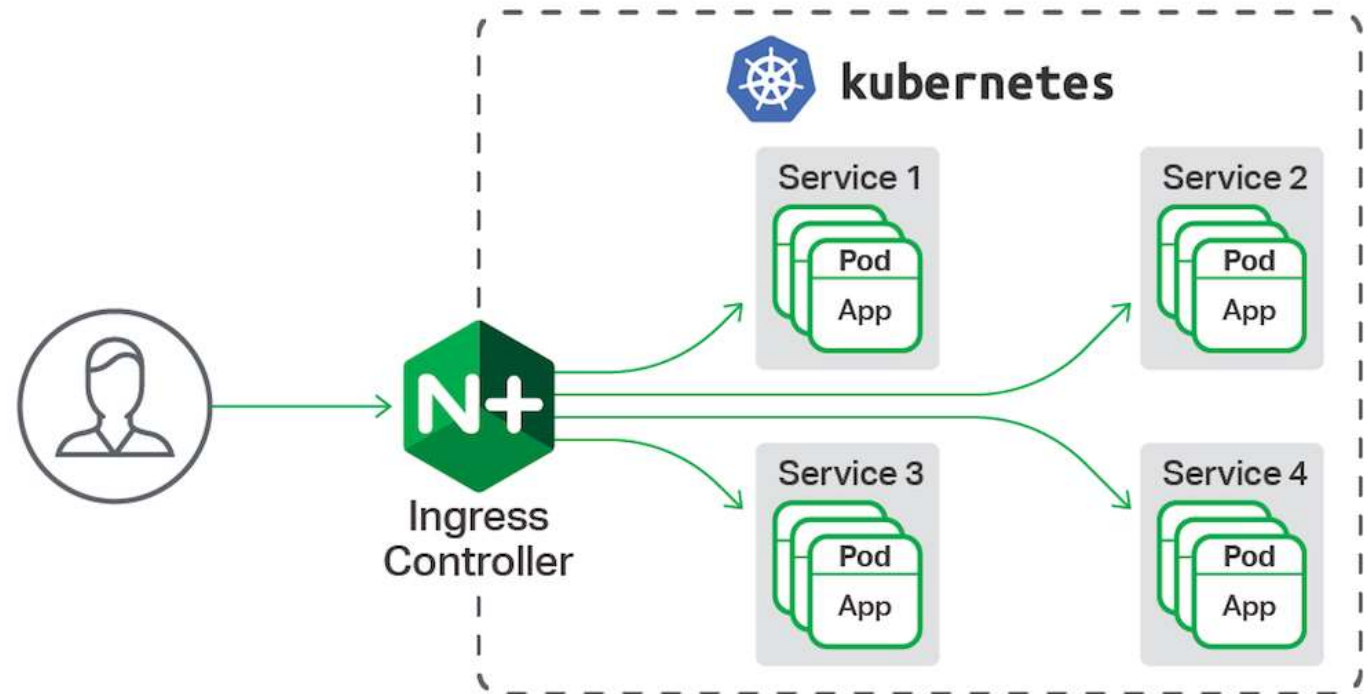
- <https://labs.play-with-k8s.com>

# Service and Ingress


- To consume your deployment, you will need to create ingress rules that expose your deployment to the external world.
- Kubernetes Ingress is a resource to add rules for routing traffic from external sources to the services in the kubernetes cluster
- To configure ingress rules in your Kubernetes cluster, first, you will need an ingress controller.
- We will create NGINX ingress controller.
  - `kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/master/deploy/static/mandatory.yaml`
  - `kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/master/deploy/static/provider/cloud-generic.yaml`

# Service and Ingress

- To confirm:
  - `kubectl get pods -n ingress-nginx`



# Service and Ingress




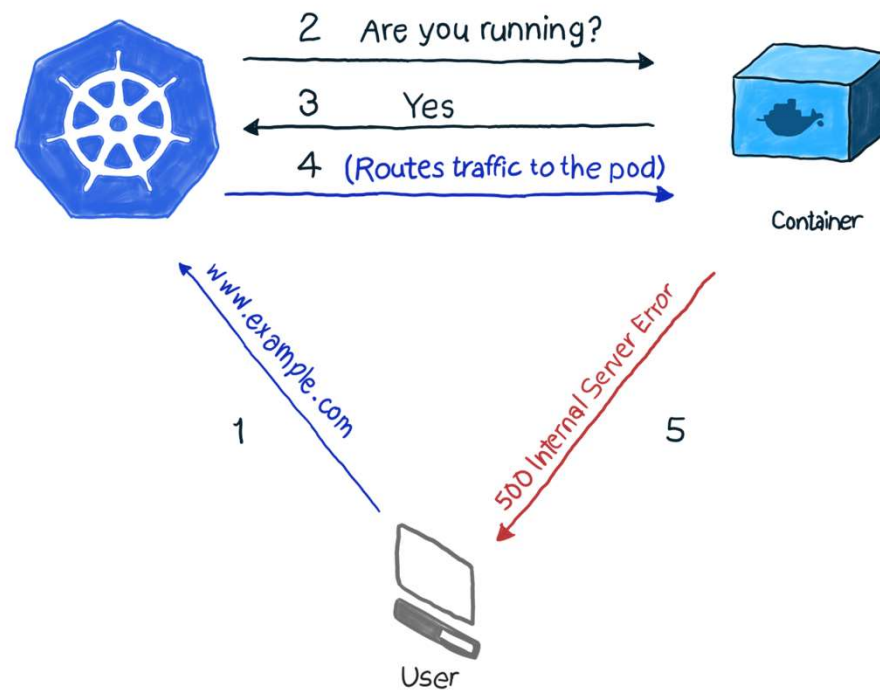
# Scaling Kubernetes

- Cluster scaling, sometimes called infrastructure-level scaling,
  - Refers to the (automated) process of adding or removing worker nodes based on cluster Utilization
- Application-level scaling, sometimes called pod scaling,
  - Refers to the (automated) process of manipulating pod characteristics based on a variety of metrics
    - CPU utilization
    - HTTP requests served per second etc
  - Two kinds of podlevel scalers exist
    - Horizontal Pod Autoscalers (HPAs), which increase or decrease the number of pod replicas depending on certain metrics.
    - Vertical Pod Autoscalers (VPAs), which increase or decrease the resource requirements of containers running in a pod.



# Scaling Kubernetes





# Thanks